

CoBERT, SPLADE-Training Retrievers

Topic: CoBERT — Contextualized Late Interaction over BERT

1. Motivation — Why CoBERT?

Earlier dense retrievers like **DPR** used **a single vector** to represent an entire passage or document.

That's efficient but **too coarse** — it loses fine-grained token-level meaning.

Example:

Passage: "Isaac Newton discovered gravity and developed calculus."

Query 1: "Who discovered gravity?"

Query 2: "Who developed calculus?"

A single embedding must represent *both ideas*, which causes ambiguity.

CoBERT was introduced to fix this by:

- Preserving **contextualized token-level embeddings** (not just one pooled vector).
- Enabling **fine-grained matching** between query tokens and document tokens.
- Maintaining **retrieval efficiency** by performing interaction **after encoding** (hence "Late Interaction").

2. Core Idea — "Late Interaction" Mechanism

Traditional BERT-based retrieval (like Cross-Encoder) encodes the **[query, passage] pair jointly**, which gives precise similarity but is **too slow** (needs re-

encoding for every query).

ColBERT introduces a compromise:

- Encode **queries and documents separately** using BERT (\rightarrow can precompute document embeddings).
- During retrieval, match them at the **token level** using a lightweight “late interaction” operation.

This gives **Cross-Encoder-level precision** with **bi-encoder-level efficiency**.

3. Architecture Overview

ColBERT has **three main components**:

◆ a. Query Encoder

- Input: a query string
- Output: contextualized token embeddings

$$Q = [q_1, q_2, \dots, q_m] \in \mathbb{R}^{m \times d}$$

(each (q_i) is a d-dimensional vector)

◆ b. Document Encoder

- Input: a passage/document
- Output: contextualized token embeddings

$$D = [d_1, d_2, \dots, d_n] \in \mathbb{R}^{n \times d}$$

◆ c. Late Interaction Layer

Instead of pooling into one vector, ColBERT computes token-level interactions:

$$\text{Sim}(Q, D) = \sum_{i=1}^m \max_{j=1, \dots, n} q_i \cdot d_j$$

That means:

- For each **query token** (q_i), find the **most similar document token** (d_j).

- Take their **dot product** (semantic similarity).
- Then sum across all query tokens for a total passage score.

This preserves token-level semantics while keeping computation scalable.



4. Training Objective — Contrastive Learning

CoBERT uses a **softmax-based contrastive loss** (similar to DPR) but with the late interaction score as similarity.

For each query (q):

- (p^+): positive (relevant) passage
- (p^-): negative (irrelevant) passages

$$L = -\log \frac{\exp(\text{Sim}(q, p^+))}{\exp(\text{Sim}(q, p^+)) + \sum_{p^-} \exp(\text{Sim}(q, p^-))}$$

This encourages the model to produce higher similarity scores for correct query-document pairs.



5. How “Late Interaction” Improves Efficiency

- **Early Interaction (Cross-Encoder):** Query and document are concatenated → huge computational cost (cannot precompute docs).
- **Late Interaction (CoBERT):** Query and document are encoded independently → documents can be pre-embedded and stored.

During retrieval, only a **lightweight dot product** is computed between token embeddings — enabling **fast ANN search** with **near-Cross-Encoder accuracy**.



6. Example — Token-Level Matching

Query: "Who discovered gravity?"

Passage: "Isaac Newton discovered gravity and developed calculus."

Step 1: Query tokens → [Who, discovered, gravity]

Step 2: Passage tokens → [Isaac, Newton, discovered, gravity, developed, calculus]

Each query token finds the *most similar* passage token:

- "Who" → "Isaac" or "Newton"
- "discovered" → "discovered"
- "gravity" → "gravity"

Then the max similarities are summed up:

$$\text{Sim} = (q_{who} \cdot d_{Newton}) + (q_{disc} \cdot d_{disc}) + (q_{grav} \cdot d_{grav})$$

✓ Result: The passage gets a high total similarity score.

7. ColBERT Workflow (Step-by-Step)

1. Indexing:

- Encode all passages using the document encoder.
- Store all their token embeddings (each doc has multiple vectors) in a **vector index** (e.g., FAISS).

2. Query Encoding:

- Encode the input query to obtain query token embeddings.

3. Retrieval:

- For each query token, retrieve the **nearest document token vectors** (ANN search).
- Combine (max + sum) similarities to score documents.
- Return top-k results.

8. Example Visualization (Text-Style)

Query: Who discovered gravity?

↓

[Who] → best match → [Newton]
[discovered] → best match → [discovered]
[gravity] → best match → [gravity]

Final Score = sum of max similarities



9. ColBERT-v2 (Improvement Highlights)

- Introduces **compressed embeddings** (from 128-dim → 32-dim)
- Uses **residual quantization** for smaller storage footprint
- Faster ANN search (up to 10× more efficient)
- Same or better retrieval accuracy

This makes ColBERT-v2 more suitable for **web-scale search** and **RAG systems**.



10. Advantages

- ✓ **High accuracy:** Close to Cross-Encoder performance
- ✓ **Fast retrieval:** Scalable via precomputed embeddings
- ✓ **Fine-grained matching:** Token-level semantic understanding
- ✓ **Interpretability:** Shows which tokens drive retrieval
- ✓ **Widely adopted:** Used in FAISS, Pyserini, and RAG pipelines



11. Limitations

- ⚠ **Storage heavy:** Each document has multiple embeddings (large vector index)
- ⚠ **Retrieval latency:** More token comparisons
- ⚠ **Complex setup:** Requires efficient vector compression and ANN indexing
- ⚠ **Training data requirement:** Needs labeled (query, passage) pairs



12. Applications

- Open-Domain QA systems** (e.g., MS MARCO, Natural Questions)
 - RAG pipelines** — precise document retrieval for LLMs
 - Search engines** — hybrid sparse+dense retrieval
 - Domain-specific retrieval** — legal, academic, or biomedical search
-



13. Key Takeaways

- **CoLBERT = Contextualized Late Interaction over BERT**
 - Introduces **token-level semantic matching** with **late interaction**
 - Balances **accuracy (like Cross-Encoder)** and **speed (like DPR)**
 - Foundation for **multi-vector retrieval models** (CoLBERT-v2, SPLADE, etc.)
 - Core building block in **retrieval-augmented generation (RAG)** systems.
-

SPLADE — Sparse Lexical and Expansion Model for Retrieval

1. Motivation — Why SPLADE?

Before SPLADE, there were two main retrieval worlds:

Type	Example	Pros	Cons
Sparse retrieval	BM25, TF-IDF	Interpretable, fast, low memory	Relies on exact word overlap
Dense retrieval	DPR, CoLBERT	Semantic understanding	Requires vector search, expensive storage

The problem:

Dense retrieval captures meaning but loses interpretability and lexical control.

Sparse retrieval is **interpretable** but misses **semantic matches**.

So researchers asked:

 Can we get the best of both — semantics of dense retrieval and explainability of sparse retrieval?

 Enter **SPLADE** — a **hybrid model** that brings neural understanding into a **sparse lexical space**.

2. Core Idea — Sparse Expansion in Vocabulary Space

SPLADE stands for:

SParse Lexical And DEnse model

But unlike DPR or ColBERT, SPLADE doesn't output continuous dense vectors.

Instead, it outputs **sparse high-dimensional vectors** aligned with the **vocabulary space** — just like TF-IDF or BM25.

How?

It uses a **transformer (like BERT)** to produce *contextualized embeddings for each token*,

then converts those embeddings into **vocabulary-level scores** using a **vocabulary projection** and **sparsifying function**.

This way:

- Each document or query is represented as a **sparse vector** over the vocabulary.
- Only a few tokens (the most meaningful ones) have **non-zero weights** — everything else is **zeroed out**.

Hence, "SPLADE" = *Sparse lexical expansion model*.

3. SPLADE Setup — Architecture and Workflow

Let's go through its setup step-by-step 

(a) Encoders

- Uses a **BERT encoder** (or any transformer) for both **queries** and **documents**.

- Each token embedding → transformed into **vocabulary logits** using a linear projection layer:

$$z_t = W \cdot h_t + b$$

where (h_t) = hidden representation of token t, and W projects to vocab size.

(b) Aggregation

- Each token predicts **weights** for many vocabulary words.
- To combine all token contributions, SPLADE uses:

$$\text{score}(w) = \max_{t \in T} (\log(1 + \text{ReLU}(z_{t,w})))$$

or sometimes a **log-sum-exp pooling** over tokens.

This creates a vector of length = vocabulary size ($\approx 30,000$ for BERT), but **most entries are zero**, giving sparsity.

(c) Sparsity Control (Regularization)

To keep the representation sparse and interpretable, SPLADE adds an **L_1 regularization term**:

$$L_{\text{sparse}} = \lambda \sum_w |\text{score}(w)|$$

This encourages many weights to be zero → sparse representation.

(d) Similarity Calculation

Once both query and document are encoded into sparse vocab-weight vectors:

$$\text{Sim}(q, d) = \sum_w q_w \cdot d_w$$

— just a **dot product** in sparse space (like BM25 cosine similarity).

 Advantage: Retrieval can be done using **traditional inverted indexes** (like Lucene or Elasticsearch).

4. Training Objective — Contrastive Learning (Like DPR)

SPLADE is trained to bring relevant query-document pairs closer and irrelevant ones farther apart.

$$L = -\log \frac{\exp(\text{Sim}(q, d^+))}{\exp(\text{Sim}(q, d^+)) + \sum_{d^-} \exp(\text{Sim}(q, d^-))}$$

and combined with sparsity loss:

$$L_{\text{total}} = L_{\text{retrieval}} + \lambda L_{\text{sparse}}$$

This balances:

- **Retrieval accuracy** (semantic matching)
 - **Sparsity** (efficiency & interpretability)
-

5. Intuition — Lexical Expansion

The “Expansion” part means SPLADE can **expand a document or query** with *semantically related words* even if they don’t appear literally.

Example:

Query: “smartphone price”

Document: “The cost of the iPhone dropped last week.”

A BM25 system would miss this because *price* ≠ *cost*.

But SPLADE learns to expand “price” → “cost”, “value”, “expense”, etc.

So in the sparse vector:

Query vector non-zeros → [smartphone, price, cost]
 Doc vector non-zeros → [iphone, cost, price]

Now, “cost” overlaps → document gets retrieved 

This expansion is **data-driven and contextual**, unlike manually crafted synonym dictionaries.

6. SPLADE Setup (Full Pipeline Summary)

Step	Component	Description
1	Transformer Encoder	Contextualizes tokens
2	Vocab Projection	Projects token embeddings → vocab dimension

Step	Component	Description
3	Pooling	Combines token vocab scores (max or log-sum-exp)
4	L ₁ Regularization	Enforces sparsity
5	Dot Product	Sparse similarity between query & doc
6	ANN/Index	Retrieval via inverted index (Lucene, Elasticsearch)

💬 7. Word Impacts in SPLADE — What Does It Mean?

“Word Impacts” refers to **which words in the vocabulary receive non-zero weights** in the SPLADE representation of a query or document.

- Each word’s weight = its **impact** (importance score)
- Words with higher weights influence retrieval more
- This impact distribution shows *which words the model thinks are important for meaning*

🧮 Mathematically:

$$\text{Impact}(w) = \log(1 + \text{ReLU}(z_w))$$

and only the top few impacts are non-zero → sparse “lexical fingerprint.”

🧠 Interpretation:

- In a **query**, high-impact words are the ones SPLADE deems essential for matching documents.
- In a **document**, high-impact words capture its main semantic themes.

Example:

Query: “What is the fastest animal on land?”

→ SPLADE might give:

```
word: impact:  
what    0
```

```
fastest 1.2
animal 1.0
land 0.8
speed 0.6 (expanded)
```

So the query representation now includes *speed*, even though it wasn't typed — this is **semantic lexical expansion**.

⚡ 8. Advantages

- ✓ Works with **existing inverted indexes** (no ANN needed).
- ✓ **Interpretable** — we can inspect word impacts.
- ✓ **Sparse but semantic** — efficient and meaningful.
- ✓ **Combines** neural understanding with lexical precision.
- ✓ High performance in **MS MARCO** and **BEIR** benchmarks.

⚠ 9. Limitations

- ⚠ Higher **training complexity** (regularization tuning).
- ⚠ Still needs **powerful transformers** → costly to train.
- ⚠ “Expansion” might introduce **noisy tokens** if not regularized well.
- ⚠ Doesn’t handle cross-modal inputs (text-only).

💡 10. Variants & Extensions

Model	Description
SPLADE (Base)	Original BERT-based sparse expansion
SPLADE-v2	Improved regularization and pooling (log-sum-exp)
DistilSPLADE	Distilled lightweight version for real-time search
SPLADE++	Hybrid fusion with dense retrievers (for RAG systems)



11. Key Takeaways

- SPLADE = Sparse Lexical and Expansion model
 - Bridges **lexical** and **semantic** retrieval.
 - Represents queries/docs as **sparse vectors** in **vocabulary space**.
 - Learns **which words to emphasize or expand** through *word impacts*.
 - Enables efficient, interpretable, semantic retrieval via **inverted indexes**.
 - Core building block for **hybrid RAG systems** (Lexical + Dense).
-

Training SPLADE Retrievers — Ranking Loss and Non-Sparsity Loss

1. Goal of Training

SPLADE is trained to **score relevant documents higher** than irrelevant ones while keeping its **representation sparse** (so that it can use inverted indexes efficiently).

So it needs to optimize **two competing objectives**:

- 1 Make good retrieval decisions → **Ranking Loss**
- 2 Keep the model efficient and interpretable → **Sparsity Regularization (Non-sparsity Loss)**

Let's break these down.

2. Ranking Loss (a.k.a. Retrieval Loss)

Intuition

When a query (q) is given, the model should give:

$\text{score}(q, d^+) > \text{score}(q, d^-)$
for all negative documents (d^-).

This is the **core ranking principle** — relevant docs should score higher than irrelevant ones.

Formulation

SPLADE uses **contrastive learning** (similar to DPR, ColBERT):

■

where:

- ($\text{sim}(q, d) = \sum_w q_w \cdot d_w$)
(dot product of sparse vectors over vocab words)
- (d^+): relevant (positive) document
- (d^-): irrelevant (negative) documents

This is essentially a **softmax ranking loss** —

encouraging the positive document's similarity to dominate over negatives.

Alternative: Margin Ranking Loss

Sometimes SPLADE variants also use:

$$L_{\text{rank}} = \max(0, m - \text{sim}(q, d^+) + \text{sim}(q, d^-))$$

where (m) is a margin hyperparameter (say, 0.2).

This pushes the positive document to be at least m more similar than any negative.

Intuitive Analogy

Think of **ranking loss** as a "teacher" telling the model:

| "For query q , doc A is good, doc B is bad — make sure A scores higher than B!"

The model learns which *vocabulary activations (word impacts)* help achieve this.

3. Non-Sparsity Loss (L_1 Regularization)

Why Needed?

Without constraint, the model might assign **non-zero scores to every vocabulary word**.

That makes retrieval:

- slow (huge index)
- uninterpretable
- memory-heavy

So we add a **non-sparsity penalty** to force the model to "focus" only on a few key words.

Formulation

$$L_{\text{sparse}} = \lambda \sum_w \text{avg}_B |\text{impact}(w)|$$

or more precisely:

$$L_{\text{sparse}} = \lambda \cdot \mathbb{E} x \in \text{batch} \sum w \log(1 + \text{ReLU}(z_{x,w}))$$

where:

- (λ) = regularization coefficient (controls sparsity strength)
 - ($z_{x,w}$) = score of vocabulary word w for query/document x
 - This penalty increases when too many vocab entries are non-zero (dense).
-

4. Combined Objective

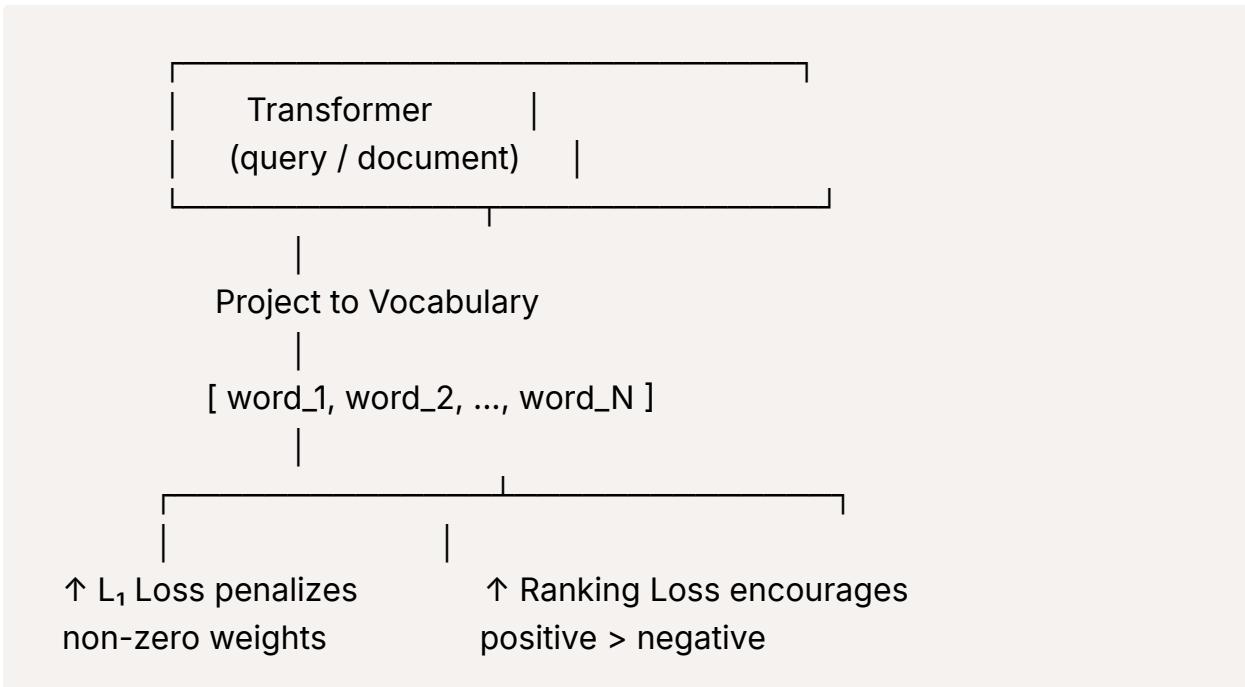
SPLADE balances both goals:

$$L_{\text{total}} = L_{\text{rank}} + \lambda L_{\text{sparse}}$$

So during training:

- **Ranking loss** → pulls relevant docs closer
 - **Sparsity loss** → pushes unimportant vocab activations toward zero
-

5. Intuitive Visualization



The L_1 loss “snips off” small weights → sparse representation.

Ranking loss tunes *which* weights survive → relevant features remain.

6. Effect of λ (Sparsity Coefficient)

λ Value	Effect
Small λ	Model keeps many non-zero tokens → less sparse but higher recall
Large λ	More sparsity → efficient, but risk of losing important terms
Balanced λ	Optimal trade-off between speed and accuracy

Tuning λ is **critical** — too much sparsity = lost meaning, too little = wasted memory.

7. Measuring Sparsity

Researchers evaluate SPLADE’s sparsity by metrics like:

- **Non-zero ratio** (fraction of active tokens in vocab)
- **Index size** (smaller = more efficient)
- **Query expansion interpretability** (are expansions meaningful?)

⚡ 8. Training Summary Table

Component	Purpose	Formula	Intuition
Ranking Loss	Make positives score higher than negatives	$(-\log \frac{e^{sim(q,d^+)}}{\sum e^{sim(q,d)}})$	Learn what "relevance" looks like
Non-Sparsity (L_1) Loss	Enforce sparsity	$\lambda \sum_w \text{avg}_B \text{impact}(w) $	pushes unimportant vocab activations toward zero
Total Loss	Combine both	$(L_{total} = L_{rank} + \lambda L_{sparse})$	Trade-off between accuracy & efficiency

Goal	What it does	Why it's important
Ranking loss	Makes relevant docs score higher than irrelevant ones	So retrieval works
Sparsity loss (L_1 loss)	Encourages fewer words to be active	So the representation stays efficient

💬 9. Example: What Happens in Practice

Query: "cheap wireless earbuds"

During training:

- Ranking loss pushes the model to make the correct docs ("affordable Bluetooth earphones") score higher.
- The model learns word expansions like:

cheap → affordable
wireless → bluetooth
earbuds → earphones

- Non-sparsity loss forces the model to **drop unhelpful activations** like:

"cheap" → ["buy", "deal", "great", ...]

if they don't add to retrieval relevance.

End result:

→ A **sparse, semantically rich, interpretable** representation.



10. Final Key Takeaways

- ✓ **Ranking Loss** teaches *what is relevant*
 - ✓ **Non-Sparsity Loss** teaches *what to ignore*
 - ✓ Together, they balance **semantic accuracy** and **index efficiency**
 - ✓ This dual-loss setup makes SPLADE's retriever both **powerful** and **practical**
-