

GANs

Discriminative Models?

A **discriminative model** is a machine learning model that **directly learns the boundary between classes** (or directly predicts y given x).

Formally:

$$P(y \mid x)$$

- They **do not model how the data is generated.**
 - Instead, they focus only on **separating classes** or predicting outputs.
-

🔑 Examples

- Logistic Regression
 - Support Vector Machines (SVMs)
 - Most Neural Networks used for classification (CNNs, RNNs, Transformers in classification tasks)
 - Conditional Random Fields (CRFs)
-

◆ How they work (intuition)

Imagine you want to classify cats vs dogs.

- A **discriminative model** looks at the features (e.g., fur length, ear shape, weight) and **draws a boundary** that best separates cats from dogs.
 - It doesn't care about *how cats look overall* or *how dogs are distributed*.
 - It only cares about **where the dividing line should be** in feature space.
-

◆ Advantages

- ✓ Often achieves **high accuracy** when lots of labeled data is available.

- ✓ **Simpler and faster** to train than generative models.
 - ✓ Flexible — you can plug in any features and learn directly.
-

◆ Limitations

- 1. Require labeled data**
 - Since they only learn $P(y | x)$, they need lots of labeled training examples.
 - They cannot easily leverage unlabeled data.
 - 2. No understanding of data distribution**
 - They don't model $P(x)$ (how the input itself is distributed).
 - So they can't **generate data, handle missing inputs gracefully, or reason about likelihood**.
 - 3. Poor generalization with limited data**
 - If data is scarce, they don't have prior knowledge of structure (unlike generative models, which can leverage assumptions about $P(x,y)$).
 - 4. Vulnerable to out-of-distribution (OOD) inputs**
 - If you give an input far from training data, they may still output a confident (but wrong) prediction.
 - Example: a cat-vs-dog classifier might confidently classify a car as "dog."
-

◆ Quick comparison

- **Generative model:** Learns full joint distribution $P(x,y)$. Can generate new samples and handle missing data. Examples: Naive Bayes, Gaussian Mixture Models, GANs.
 - **Discriminative model:** Learns only $P(y | x)$. Great at classification, but nothing more.
-

- ✓ **Summary in one line:**

Discriminative models directly learn decision boundaries ($P(y|x)$) and are great for classification, but they can't model the data distribution, struggle with limited

labeled data, and fail on unseen/out-of-distribution inputs.

Generative Models

A **generative model** tries to learn the **joint probability distribution**:

$$P(x, y) = P(x | y)P(y)$$

- It models **how the data is generated**:
 1. Pick a label y .
 2. Generate input x conditioned on y .
- Once we know $P(x,y)$, we can derive the classifier using Bayes' rule:

$$P(y | x) = \frac{P(x|y)P(y)}{P(x)}$$

◆ How do they cope with the Discriminative model limitations

1. Limited labeled data

- Since they learn **how data is distributed** ($P(x | y)$), they can leverage **unlabeled data** to improve learning.
 - Example: Semi-supervised generative models can use unlabeled images to learn $P(x)$, then use a small amount of labels for $P(y | x)$.
-

2. Understanding the data distribution

- Generative models explicitly model $P(x)$.
 - This means they **know what valid data looks like**.
 - Applications:
 - Detecting **out-of-distribution** samples (OOD detection).
 - **Denoising** or imputing missing data.
 - Generating synthetic samples (new data).
-

3. Better generalization

- By capturing the structure of data, they can generalize better when labeled data is scarce.
 - Example: A Gaussian Naive Bayes classifier can perform well even with few training points, because it assumes a generative structure (features conditioned on class follow Gaussians).
-

4. Out-of-distribution robustness

- Since they know what the data distribution $P(x)$ looks like, they can recognize when an input **doesn't belong** to the training distribution.
 - Example: A discriminative cat-vs-dog classifier might label a car as "dog," while a generative model can say, "This sample doesn't look like cats or dogs I know."
-

◆ Examples of Generative Models

- **Classic (probabilistic):**
 - Naive Bayes
 - Gaussian Mixture Models (GMM)
 - Hidden Markov Models (HMM)
 - **Modern (deep learning):**
 - Variational Autoencoders (VAEs)
 - Generative Adversarial Networks (GANs)
 - Diffusion Models
-

◆ Tradeoff

- **Discriminative models** → better at classification when you have *lots of labeled data*.
 - **Generative models** → more flexible, can work with less labeled data, handle OOD detection, and can generate data... but often harder to train and less accurate in pure classification tasks.
-

Summary in one line:

Generative models overcome the main weaknesses of discriminative models by learning the **data distribution itself** ($P(x,y)$), which lets them use unlabeled data, detect outliers, generate samples, and generalize better with scarce labels.

Implicit Density Estimation Problem

1. Background: Density Estimation

- In probability and machine learning, **density estimation** means trying to learn the probability distribution $p(x)$ that generates the data.
- Example: if we have handwritten digit images, we'd like to know the probability distribution over all possible digit images.

There are two broad ways to do this:

- **Explicit Density Estimation:** We define a probability model with a tractable form and learn its parameters.
 - Example: Gaussian Mixture Models (GMMs), Autoregressive models, Normalizing Flows.
 - They allow exact evaluation of $p(x)$.
 - **Implicit Density Estimation:** We don't explicitly define or compute $p(x)$. Instead, we **learn a model that can sample from the distribution**, even if we can't write down its probability function.
-

2. What is the Implicit Density Estimation Problem?

- In **implicit models** (like GANs), we can generate samples that *look like* the data.
- However, we **don't have direct access to the probability density function $p(x)$** .
- That means:
 - We **can't evaluate** how likely a given point is under the learned model.
 - We **can't compute likelihood-based metrics** directly.

- Training becomes harder because we can't just maximize log-likelihood.
-

3. Example: GANs

- GANs define a generator $G(z)$ that transforms noise $z \sim p(z)$ into samples $x = G(z)$.
 - This gives us a way to sample from $p_\theta(x)$, but we **can't compute the probability density** of any particular x .
 - The discriminator helps indirectly guide the generator without computing likelihoods.
-

4. Why is it a "Problem"?

- Because without explicit density:
 - It's harder to measure how good the model is (no exact log-likelihood).
 - It's harder to combine with probabilistic reasoning tasks (e.g., anomaly detection).
 - Training relies on surrogate losses (e.g., adversarial loss, divergence minimization).
-

Key takeaway:

- **Explicit models** = "I can *describe* the probability distribution."
 - **Implicit models** = "I can't describe it, but I can *mimic* it and generate samples from it."
-

Summary in Simple Words:

The **implicit density estimation problem** is that some generative models (like GANs) can generate data that looks realistic, but **cannot tell you the actual probability of that data**. They only learn to **mimic** the distribution, not to describe it mathematically.

Generative Adversarial Networks (GANs)

1. Core Idea

GANs introduce a clever *game* between two neural networks:

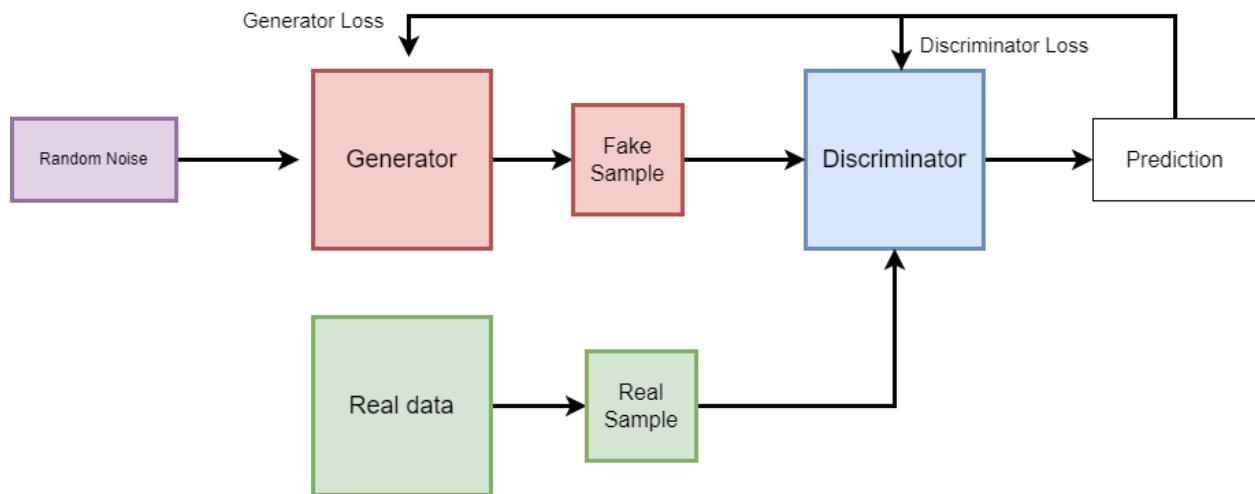
- **Generator (G)**: Learns to produce fake samples that resemble real data.
- **Discriminator (D)**: Learns to distinguish between real data and fake (generated) data.

They are trained together in a **minimax game**:

- G tries to fool D.
- D tries to catch G.

This competition pushes the generator to produce increasingly realistic samples.

2. Architecture



1. Generator (G)

The **generator's job** is to take in random noise and transform it into realistic-looking data (images, audio, text, etc.).

◆ Input

- A random vector $z \sim p(z)$.
- Typical choice: Gaussian $\mathcal{N}(0, I)$ or Uniform distribution.

- Dimension: usually much smaller than the output dimension (e.g., $z \in \mathbb{R}^{100}$ for an image of size $64 \times 64 \times 3$).

◆ Layers

- **Fully connected layers (MLPs)** → map noise to a higher-dimensional representation.
- **Convolutional layers (in image GANs):**
 - Use **transposed convolution (a.k.a. deconvolution)** to progressively upsample from a small feature map to a large image.
 - Example: start from a 4×4 feature map → upsample to $8 \times 8 \rightarrow 16 \times 16 \rightarrow 64 \times 64$.
- **Batch Normalization** → stabilizes training.
- **Activation functions:**
 - ReLU or LeakyReLU in hidden layers.
 - **Tanh** in the final layer (so pixel values fall in $[-1,1]$).

◆ Output

- A fake data sample $x_{fake} = G(z)$.
 - Example: $64 \times 64 \times 3$ RGB image.
-

2. Discriminator (D)

The **discriminator's job** is to classify whether input data is real (from dataset) or fake (from generator).

◆ Input

- A data sample x (either real or fake).
- Example: $64 \times 64 \times 3$ image.

◆ Layers

- **Convolutional layers (for images):**
 - Reduce spatial size (downsampling) while increasing depth.

- Example: $64 \times 64 \times 3 \rightarrow 32 \times 32 \times 64 \rightarrow 16 \times 16 \times 128$.
- **Batch Normalization** → sometimes avoided in the first layer to avoid gradient issues.
- **Activation functions:**
 - LeakyReLU (to prevent dying ReLU problem).

◆ Output

- A single probability value $D(x) \in [0,1]$.
- Represents the likelihood that x is real.
- Uses **Sigmoid activation** at the last layer.

3. Training Interaction

GAN training is a **two-player minimax game**:

- The **Discriminator (D)** tries to classify samples correctly (real = 1, fake = 0).
- The **Generator (G)** tries to fool D (make fake samples look real).

1. Discriminator Loss

The discriminator's job:

- Maximize the probability of classifying real data as real.
- Maximize the probability of classifying fake data as fake.

Mathematically:

$$L_D = - \left(\mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))] \right)$$

- First term: reward D for outputting high probability on **real data**.
- Second term: reward D for outputting low probability on **fake data**.

👉 Training update: **minimize L_D** w.r.t. discriminator's weights.

(Equivalently, maximize the log-likelihood of correct classification.)

2. Generator Loss

The generator's job:

- Fool the discriminator into believing generated samples are real.

Two versions exist:

(a) **Original GAN Loss (Minimax)**

$$L_G = \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

- Generator minimizes this.
- But problem: if D is too strong early, $D(G(z)) \approx 0$, log saturates \rightarrow **vanishing gradients**.

(b) **Non-saturating Loss (Practical Version)**

$$L_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

- Equivalent to maximizing $\log D(G(z))$.
- Gives stronger gradients when D is good.
- This is the **standard choice in practice**.

👉 Training update: minimize L_G w.r.t. generator's weights.

3. Training Algorithm (High-level)

1. Sample a minibatch of **real data** $x \sim p_{data}(x)$.
2. Sample random noise $z \sim p(z)$.
3. Generate fake samples $G(z)$.
4. Update **Discriminator**:
 - Minimize L_D (real $\rightarrow 1$, fake $\rightarrow 0$).
5. Update **Generator**:
 - Minimize L_G (make D think fake is real).
6. Repeat steps until convergence (or until samples look realistic).

\sim : This symbol is called a "tilde". In this context, it means "**is drawn from**" or "**is distributed as**". It indicates that the variable on the left is a random sample from the probability distribution on the right.

4. Intuition Recap

- **Discriminator Loss** = "How well am I at telling real vs fake?"
- **Generator Loss** = "How well am I at fooling D into thinking fake = real?"
- Training is a **tug-of-war**:
 - If **D is too strong**, G can't learn (no gradient).
 - If **G is too strong**, D becomes useless.
 - Balance is key → that's why GAN training is tricky.

5. Summary:

- **Discriminator Loss**:

$$L_D = -[\log D(x) + \log(1 - D(G(z)))]$$

- **Generator Loss** (practical version):

$$L_G = -\log D(G(z))$$

4. Data Flow Summary

1. Noise $z \rightarrow$ Generator $G(z) \rightarrow$ Fake sample x_{fake} .
2. Real sample x or fake sample $x_{fake} \rightarrow$ Discriminator $D(x)$.
3. D outputs probability of real vs fake.
4. Losses update both networks in opposite directions.

5. Architecture Variants

- **Vanilla GAN (MLPs only)** → early versions.

- **DCGAN (Deep Convolutional GAN)** → uses conv & deconv layers for images, much more powerful.
 - **WGAN (Wasserstein GAN)** → stabilizes training using Wasserstein distance.
 - **Conditional GAN (cGAN)** → conditions G and D on labels (e.g., generate a cat vs a dog).
 - **StyleGAN** → advanced architecture for photorealistic faces.
-

Intuition Recap:

- Generator = **artist** (paints fake images from imagination).
 - Discriminator = **art critic** (judges whether the painting is real or fake).
 - Training = **competition** that makes the artist improve until the critic can't tell real from fake.
-

3. Objective Function

The GAN loss is a **minimax optimization**:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

- **Discriminator maximizes:** classify real as 1, fake as 0.
 - **Generator minimizes:** fool D (make $D(G(z)) \approx 1$).
-

4. Training Dynamics

- Step 1: Train **D** to better separate real vs fake.
 - Step 2: Train **G** so that its fake samples get better at fooling D.
 - Alternate between D and G updates.
 - Ideally, the system converges when $p_{model}(x) = p_{data}(x)$ (generator distribution matches real distribution).
-

5. Advantages of GANs

- ✓ Can generate very realistic samples (images, music, text).
 - ✓ No need for explicit likelihood computation (solves the implicit density problem).
 - ✓ Highly expressive since G is a neural net.
-

6. Limitations of GANs

- ▲ Training instability (non-convex minimax optimization).
 - ▲ Mode collapse: G produces limited diversity (e.g., always generates similar faces).
 - ▲ Hard to evaluate progress quantitatively (no likelihood).
-

7. Real-world Applications

- Image synthesis (e.g., **DeepFake**, StyleGAN).
 - Text-to-image models (when combined with transformers).
 - Data augmentation.
 - Super-resolution (making images sharper).
 - Art and creative content generation.
-

✓ Intuition Summary:

GANs are like a **counterfeit money maker (Generator)** competing with a **police detective (Discriminator)**.

- The generator wants to make fake currency that looks real.
 - The discriminator wants to spot fakes.
 - Over time, both get better → the generator eventually produces “fakes” that are indistinguishable from real data.
-