

# Language Modeling

---

## Language Modeling

### Intuitive Understanding

**Language Modeling** is about teaching a machine to **understand and generate human language** — by learning the **probability of word sequences**.

Think of it as helping the model answer:

| "Given the words so far, what word (or token) is most likely to come next?"

### Example

Suppose you have these two sentences:

1. "The cat sat on the \_\_\_\_."
2. "The cat sat on the airplane."

A **language model** will assign a **higher probability** to:

| "The cat sat on the mat."

because "mat" is a more likely continuation based on how words co-occur in real language.

### Formal Definition

A **Language Model (LM)** estimates the **probability distribution** over sequences of words (or tokens):

$$P(w_1, w_2, w_3, \dots, w_T)$$

For practical use, it learns to compute:

$$P(w_t | w_1, w_2, \dots, w_{t-1})$$

This means:

| The probability of the current word w<sub>t</sub>w<sub>t</sub>w<sub>t</sub> given all previous words.

---

## Types of Language Models

### 1. Statistical (Classical) Language Models

Before deep learning, LMs used counting-based approaches.

- **Unigram model:** assumes all words are independent  
→  $P(w_1, w_2, \dots, w_T) = \prod_i P(w_i)$
- **Bigram model:** depends only on previous one word  
→  $P(w_t | w_{t-1})$
- **Trigram model:** depends on last two words  
→  $P(w_t | w_{t-2}, w_{t-1})$

➡ **Limitation:** Can't handle long-term dependencies well.

---

### 2. Neural Language Models

Instead of counting, they **learn** patterns using neural networks.

#### a. Feedforward Neural LMs

- Inputs are fixed-size windows of previous words.
- Output: probability distribution over next word.

#### b. Recurrent Neural Networks (RNN / LSTM / GRU)

- Maintain a hidden state capturing past context.
- Can process sequences of any length.

#### c. Transformer-based Models (Modern)

- Use **self-attention** to capture dependencies between all words in a sequence simultaneously.
- Example models:
  - **GPT (Generative Pretrained Transformer)** — autoregressive LM

- **BERT** — masked LM (predict missing words)
- 

# How to build a neural language model?

1. **Gather text data** (corpus) — e.g., WikiText, books, scraped text.
  2. **Tokenize** — map text → tokens. Use subword (BPE / WordPiece / SentencePiece) for open vocab.
  3. **Create dataset** — convert token ids into training sequences (inputs and targets). Use sliding windows or next-token pairs.
  4. **Model** — choose architecture: RNN/LSTM/GRU (simple), Transformer (state of the art).
  5. **Loss** — Cross-entropy on next-token prediction.
  6. **Optimization** — Adam/AdamW, LR schedule (warmup + decay), gradient clipping.
  7. **Evaluation** — Perplexity ( $\exp(\text{avg cross-entropy})$ ), and sample quality.
  8. **Inference** — greedy, beam, top-k, top-p (nucleus) sampling for generation.
  9. **Deploy/Serve** — convert to ONNX / TorchScript or use model-serving infra.
- 

# Fixed-Window Neural Language Model

## 1. Intuitive Idea

A **Fixed-window Neural Language Model (NNLM)** predicts the next word based on a **fixed number of previous words** (just like an  $n$ -gram),

But instead of counting co-occurrences, it **learns distributed word embeddings** and uses a **neural network** to generalize.

Think: "Instead of memorizing all possible 3-word combinations, I'll learn continuous representations of words that let me guess likely next words, even for unseen sequences."

---

## 2. The Setting

Let's say you have a sequence of words:

$$w_1, w_2, w_3, \dots, w_T$$

We want to estimate:

$$P(w_t | w_{t-n+1}, \dots, w_{t-1})$$

This is similar to a trigram or 5-gram model — we only look at a *fixed window* of  $n-1$  previous words.

---

## 3. Mathematically

Let's define:

- $V$ : vocabulary size
- $d$ : embedding dimension
- $n$ : context window size (number of previous words used)

Each word  $w_i$  is represented as a one-hot vector  $x_i \in \mathbb{R}^V$ .

### Step 1: Embedding lookup

We learn an embedding matrix  $E \in \mathbb{R}^{V \times d}$ .

$$e_i = E^T x_i$$

Now we have  $n-1$  embeddings:  $e_{t-n+1}, \dots, e_{t-1}$ .

### Step 2: Concatenate

$$z = [e_{t-n+1}; e_{t-n+2}; \dots; e_{t-1}] \in \mathbb{R}^{(n-1)d}$$

### Step 3: Feedforward neural network

We apply a non-linear transformation:

$$h = \tanh(W_1 z + b_1)$$

Then compute scores for all words in vocabulary:

$$o = W_2 h + b_2$$

## Step 4: Softmax output

$$P(w_t = i | context) = \frac{\exp(o_i)}{\sum_j \exp(o_j)}$$

## Step 5: Training

Use **cross-entropy loss** to maximize the probability of the correct next word.

---

## 4. What It Learns

- The **embedding matrix E** learns semantic representations of words (similar words get similar vectors).
  - The **neural layers** learn to **combine** context information nonlinearly.
  - This overcomes data sparsity of n-grams: it can generalize from *similar* contexts.
- 

## 5. Limitations

Problem	Why It Matters
<b>Fixed window</b>	Can't use longer context (e.g., 10+ previous words) without huge parameter growth.
<b>Parameter explosion</b>	Input layer grows linearly with window size $\times$ embedding dim.
<b>No sequence memory</b>	Doesn't "remember" beyond window — no notion of sentence history.

That's why **RNNs** and later **Transformers** replaced it — they can handle **variable-length** context and **capture long dependencies**.