

# Popular Optimizers in Deep Learning (GD, SGD, Adam, etc.)

---

## High-Level Summary

Optimizers are algorithms that **update the model's parameters** (weights, biases) to minimize the loss.

They differ in **speed, stability, memory use, and convergence behavior**.

---

## Detailed Explanation

### ◆ 1. Gradient Descent (Batch GD)

- **What it is:** Uses the **entire dataset** to compute gradients before updating weights.
- **Update rule:**

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

- **Why used:** Simple, stable for convex functions.
  - **Limitations:** Very slow for large datasets (one update per full dataset pass).
  - **When to use:** Small datasets (fits in memory). Rare in modern deep learning.
- 

### ◆ 2. Stochastic Gradient Descent (SGD)

- **What it is:** Updates parameters using **one sample at a time**.
- **Update rule:**

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(x_i, y_i; \theta)$$

- **Why used:** Much faster than GD, introduces **noise** that helps escape local minima.

- **Limitations:** Very noisy, loss oscillates.
  - **When to use:** Streaming data, online learning.
- 

### ◆ 3. Mini-Batch SGD (default in practice)

- **What it is:** Updates using a **small batch** of samples (e.g., 32, 64).
  - **Why used:**
    - More efficient than GD.
    - Smoother than pure SGD.
    - Well-suited for GPUs.
  - **When to use:** Standard choice for deep learning training.
- 

### ◆ 4. SGD with Momentum

- **What it is:** Adds a **velocity term** to smooth updates, accumulates past gradients.
- **Update rules:**

$$v_t = \beta v_{t-1} + (1 - \beta) \nabla_{\theta} \mathcal{L}$$

$$\theta \leftarrow \theta - \eta v_t$$

- **Why used:** Speeds up convergence, prevents oscillations.
  - **When to use:** Deep CNNs (ResNet, VGG).
- 

### ◆ 5. RMSProp

- **What it is:** Adapts learning rate per parameter by dividing by a moving average of squared gradients.
- **Update rule:**

$$s_t = \beta s_{t-1} + (1 - \beta) (\nabla_{\theta} \mathcal{L})^2$$

— —

$$\theta \leftarrow \theta - \eta \frac{\nabla_{\theta} \mathcal{L}}{\sqrt{s_t + \epsilon}}$$

- **Why used:** Handles non-stationary problems well, stabilizes learning.
  - **When to use:** RNNs, unstable training tasks.
- 

## ◆ 6. Adam (Adaptive Moment Estimation)

- **What it is:** Combines **Momentum** + **RMSProp** (first + second moment estimates).
- **Update rules:**

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla_{\theta} \mathcal{L}$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla_{\theta} \mathcal{L})^2$$

Bias correction:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

Update:

$$\theta \leftarrow \theta - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$$

- **Why used:**
    - Fast convergence.
    - Works well with little tuning.
    - Handles sparse gradients.
  - **When to use:** Default for NLP, Transformers, GANs.
- 

## ◆ 7. AdamW

- **What it is:** Adam with **decoupled weight decay** (fixes over-regularization issue).

- **Why used:** More stable than Adam, better generalization.
  - **When to use:** Transformers (BERT, GPT, ViTs).
- 

## Analogy

- **GD:** Like checking the **entire class's exams** before adjusting teaching.
  - **SGD:** Like checking **one student's exam** after each question.
  - **Mini-batch SGD:** Like checking a **small group** of exams before changing teaching.
  - **Momentum:** Like a ball rolling downhill — builds speed in the right direction.
  - **RMSProp:** Like adjusting your stride size based on the terrain.
  - **Adam:** Combines momentum (rolling ball) + adaptive stride (terrain-aware).
- 

## Comparison Table

Optimizer	Idea	Pros	Cons	When to Use
<b>GD</b>	Full dataset update	Stable, exact	Slow, memory heavy	Small datasets
<b>SGD</b>	One sample update	Fast, helps escape minima	Noisy updates	Online learning
<b>Mini-batch SGD</b>	Small batch	Efficient, stable	Needs batch tuning	Deep learning default
<b>Momentum SGD</b>	Adds velocity	Faster convergence	Needs momentum tuning	Deep CNNs
<b>RMSProp</b>	Scales lr by past gradients	Handles exploding/vanishing gradients	May generalize poorly	RNNs
<b>Adam</b>	Momentum + RMSProp	Fast, works out of box	May overfit, bad generalization	NLP, GANs
<b>AdamW</b>	Adam + weight decay	Better generalization	Slightly slower	Transformers, SOTA models

## Key Takeaway

- Use **Mini-batch SGD + Momentum** for CNNs.
  - Use **Adam/AdamW** for Transformers, NLP, GANs.
  - Use **RMSProp** for RNNs or unstable cases.
  - Use **GD** only for small datasets.
-