

Training Neural Networks

Loss Functions

High-Level Summary

A **loss function** is a mathematical formula that measures **how wrong a model's prediction is** compared to the true label.

It provides the **signal for learning** by guiding how the model's parameters should be updated during training.

Detailed Explanation

When a model makes predictions, we need a way to **quantify the error**.

- **Loss function:** Calculates error for a single prediction.
- **Cost function:** Usually refers to the average loss across all training examples.

Loss functions differ depending on the **type of task**:

- ◆ **1. Regression Loss Functions (continuous outputs)**
- **Mean Squared Error (MSE):**

$$L = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- Penalizes larger errors more heavily.
- Common in predicting prices, temperatures, etc.
- **Mean Absolute Error (MAE):**

$$L = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

- Robust to outliers, treats all errors equally.
 - **Huber Loss:**
 - Combination of MSE + MAE (less sensitive to outliers than MSE).
-

◆ **2. Classification Loss Functions (discrete outputs)**

- **Binary Cross-Entropy (Log Loss)** (for binary classification):

$$L = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

- Used with **Sigmoid** in binary tasks.
- Example: spam vs. not spam.
- **Categorical Cross-Entropy** (for multi-class classification):

$$L = - \sum_{i=1}^N \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c})$$

- Used with **Softmax**.
 - Example: image classification (dog, cat, car).
-
- ◆ **3. Specialized Loss Functions**
- **Hinge Loss:** Used in Support Vector Machines (SVMs).
 - **KL Divergence:** Measures difference between probability distributions.
 - **IoU (Intersection over Union) Loss / Dice Loss:** Used in **segmentation** tasks.
 - **Adversarial Loss:** Used in GANs (generator vs. discriminator).
 - **Contrastive Loss / Triplet Loss:** Used in **metric learning** (face verification, embeddings).
-

Analogy

Think of the loss function like a **teacher's red pen**:

- Every time you answer (predict), the teacher marks how wrong you are.
 - The goal is to keep adjusting (learning) until the red marks (loss) are minimized.
-

Mathematical Foundation

General idea:

$$(y, \hat{y}) = \text{Error between true value } y \text{ and prediction } \hat{y}$$

Training goal:

$$\theta^* = \arg \min_{\theta} \frac{1}{N} \sum_{i=1}^N L(y_i, f_{\theta}(x_i))$$

Where:

- f_{θ} = model with parameters θ .
 - y_i = true labels.
 - \hat{y}_i = predictions.
-

Use Case Examples

- Predicting **house prices** → MSE (regression).
 - Detecting **spam emails** → Binary Cross-Entropy.
 - Classifying **ImageNet images** → Categorical Cross-Entropy.
 - Segmenting **medical images** → Dice Loss / IoU Loss.
 - Training **GANs** → Adversarial Loss.
-

✓ Key Insight:

- **Regression** → MSE / MAE

- **Classification** → **Cross-Entropy**
 - **Segmentation** → **Dice / IoU**
 - **Generative / Embedding tasks** → **Specialized losses**
-

Convex Loss Function

- In optimization, we want to **minimize the loss function** to find the best model parameters.
 - If the loss function is **convex**, gradient descent is guaranteed (in theory) to reach the **global minimum**.
 - If it is **non-convex**, there may be many local minima/saddle points → optimization becomes harder (common in deep learning).
 - A **convex loss function** is one where the error curve is shaped like a **bowl (U-shape)** — meaning it has **one global minimum** and no local minima, making optimization easier and more stable.
-

Examples of Convex Loss Functions

1. **Mean Squared Error (MSE)**
 - Quadratic → convex.
 - Common in regression.
 2. **Hinge Loss** (used in SVMs)
 - Convex because it is piecewise linear.
 3. **Log Loss / Cross-Entropy** (for classification)
 - Convex with respect to predictions.
-

Gradient Descent & Stochastic Gradient Descent

High-Level Summary

- **Gradient Descent (GD)** is an optimization method that updates model parameters in the **direction of the negative gradient** of the loss to minimize it.

- **Stochastic Gradient Descent (SGD)** is a faster variant that updates parameters using **one (or a few) training examples at a time**, introducing randomness that helps escape poor local minima.
-

Detailed Explanation

- ◆ **1. Gradient Descent (Batch Gradient Descent)**
 - Compute the loss over the **entire dataset**.
 - Compute the gradient of loss w.r.t. each parameter.
 - Update all parameters **once per pass (epoch)**.

Update rule:

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

Where:

- θ = parameters (weights, biases)
- η = learning rate (step size)
- $\nabla_{\theta} \mathcal{L}(\theta)$ = gradient of loss function

Pros: Stable, guaranteed convergence for convex losses.

Cons: Very **slow** for large datasets because each update requires scanning all examples.

- ◆ **2. Stochastic Gradient Descent (SGD)**

- Instead of using all data, update parameters after **just one sample** (or a small random subset).
- Introduces **stochasticity (noise)** into updates.

Update rule (per sample):

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(x_i, y_i; \theta)$$

Pros:

- Much faster for large datasets.
- Noise helps **escape saddle points and local minima**.

Cons:

- Updates are noisy, loss may oscillate.
 - Convergence less stable without tricks like learning rate schedules or momentum.
-

◆ 3. Mini-Batch Gradient Descent (Hybrid)

- Uses a **small batch** of data (e.g., 32, 64, 128 samples).
- Compromise between **efficiency** and **stability**.
- Modern deep learning almost always uses **mini-batch SGD**.

Update rule (per mini-batch):

$$\theta \leftarrow \theta - \eta \frac{1}{m} \sum_{j=1}^m \nabla_{\theta} \mathcal{L}(x_j, y_j; \theta)$$

Where m = batch size.

Analogy

- **Batch Gradient Descent:** Like checking the **entire class's exam papers** before deciding how to improve teaching.
 - **SGD:** Like checking **just one student's paper** and immediately changing teaching style.
 - **Mini-batch GD:** Like checking a **small group's papers** (say 10 students) before making adjustments.
-

Mathematical Foundation

1. We want to solve:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta)$$

1. Gradient descent updates iteratively:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} \mathcal{L}(\theta_t)$$

1. For mini-batches:

$$\nabla_{\theta} \mathcal{L}(\theta_t) \approx \frac{1}{m} \sum_{j=1}^m \nabla_{\theta} \mathcal{L}(x_j, y_j; \theta_t)$$

Use Cases

- **Batch GD:** Small datasets where full dataset fits in memory (e.g., linear regression on small data).
 - **SGD:** Online learning, streaming data, very large datasets.
 - **Mini-batch GD:** Standard choice for deep learning (vision, NLP, speech).
-

Extra Insight

- SGD is often paired with **optimizers** that improve convergence:
 - **Momentum:** accelerates in consistent gradient directions.
 - **Adam, RMSProp:** adaptive learning rates.
 - In practice, when people say “**SGD**” in deep learning, they almost always mean **mini-batch SGD**.
-

✓ Key Takeaway:

- **Batch GD** = exact but slow.
 - **SGD** = noisy but fast.
 - **Mini-batch SGD** = sweet spot → standard in deep learning.
-

Vanishing & Exploding Gradients

High-Level Summary

- The **Vanishing Gradient Problem** happens when gradients become extremely **small** during backpropagation, making early layers learn very slowly (or stop learning).
 - The **Exploding Gradient Problem** happens when gradients become extremely **large**, causing unstable updates and diverging weights.
-

Detailed Explanation

◆ 1. Why These Problems Occur

- In deep networks, backpropagation uses the **chain rule**:

$$\frac{\partial \mathcal{L}}{\partial W^{[1]}} = \frac{\partial \mathcal{L}}{\partial a^{[L]}} \cdot \frac{\partial a^{[L]}}{\partial z^{[L]}} \cdots \frac{\partial a^{[2]}}{\partial z^{[2]}} \cdot \frac{\partial z^{[2]}}{\partial W^{[1]}}$$

- This is a **product of many derivatives**.
 - If derivatives $< 1 \rightarrow$ product shrinks \rightarrow **vanishing gradient**.
 - If derivatives $> 1 \rightarrow$ product explodes \rightarrow **exploding gradient**.
-

◆ 2. Vanishing Gradient

- Common with **sigmoid** or **tanh** activations, where:
 - $\sigma'(z) = \sigma(z)(1 - \sigma(z)) \leq 0.25$.
 - $\tanh'(z) \leq 1$.
- In deep nets, multiplying numbers < 1 repeatedly \rightarrow gradient goes to **0**.

Effect:

- Early layers don't update \rightarrow network can't learn long dependencies.
 - RNNs suffered heavily from this problem before LSTM/GRU.
-

◆ 3. Exploding Gradient

- Opposite effect: derivatives (or weights) are large.
- Multiplying numbers > 1 repeatedly \rightarrow gradient grows **exponentially**.

Effect:

- Weight updates become huge.
 - Loss oscillates or goes to **NaN** (not a number).
-

◆ 4. Causes

- Poor weight initialization (too large/small).

- Very deep architectures.
 - Certain activation functions (sigmoid, tanh are more prone).
 - Recurrent networks (due to repeated multiplications over time steps).
-

Analogy

- **Vanishing Gradient:** Like whispering a message down a **long hallway**. By the time it reaches the first person, the signal is so faint they can't hear it.
 - **Exploding Gradient:** Like shouting through a **microphone with max volume**. The signal becomes so loud that it distorts everything.
-

Mathematical Foundation

Suppose each layer multiplies by a Jacobian matrix $J[l]J^{\wedge\{[l]\}}$.

The gradient is:

$$\frac{\partial \mathcal{L}}{\partial x} = \prod_{l=1}^L J^{[l]}$$

- If eigenvalues of $J^{[l]}$ are < 1 , the product tends to $\mathbf{0}$ \rightarrow vanishing gradient.
 - If eigenvalues are > 1 , the product tends to ∞ \rightarrow exploding gradient.
-

Solutions

For Vanishing Gradient

- Use **ReLU** (derivative is 0 or 1, not shrinking like sigmoid).
- Use **Batch Normalization**.
- Use **Residual Connections (ResNets)**.
- Use **LSTMs/GRUs** in RNNs (gating helps preserve gradient flow).
- Careful weight initialization (Xavier/He init).

For Exploding Gradient

- **Gradient Clipping** (cap gradients at a threshold).
- Smaller learning rates.

- Careful initialization.
-

Use Case

- In training **deep CNNs**, vanishing gradients can stall learning in early layers → ResNet skip connections solved this.
 - In **RNNs**, exploding gradients caused instability → gradient clipping + gated RNNs fixed it.
-

✓ Key Takeaway:

- **Vanishing** = network can't learn early features.
 - **Exploding** = network updates blow up, training diverges.
-

Backpropagation (Forward & Backward Propagation)

High-Level Summary

Backpropagation is the algorithm that **efficiently computes gradients** of a loss with respect to all network parameters by applying the **chain rule** backward through the network; the **forward pass** computes predictions and loss, the **backward pass** propagates error signals to obtain gradients used to update weights.

Detailed Explanation

1) Forward Propagation (compute predictions & loss)

For an L -layer feedforward network (MLP/CNN head notation):

- Inputs: $a^{[0]} = x$
- For each layer $l = 1, \dots, L$:
 - **Affine/conv step:** $z^{[l]} = W^{[l]}a^{[l-1]} + b^{[l]}$
 - (For CNNs, $W^{[l]}$ applies a convolution instead of a matrix multiply.)

- **Nonlinearity:** $a^{[l]} = \phi^{[l]}(z^{[l]})$
- **Loss** (example: cross-entropy for classification): $\mathcal{L} = \ell(a^{[L]}, y)$

You also **cache** intermediates $\{z^{[l]}, a^{[l]}\}$ for the backward pass.

2) Backward Propagation (compute gradients)

Start from the loss and move **right → left** (output → input).

Define the **error signal** (a.k.a. **delta**) at layer l :

$$\delta^{[l]} \triangleq \frac{\partial \mathcal{L}}{\partial z^{[l]}}$$

Then, for $l = L, L-1, \dots, 1$:

1. Local gradient at layer output

$$\delta^{[l]} = \frac{\partial \mathcal{L}}{\partial a^{[l]}} \odot \phi'^{[l]}(z^{[l]})$$

where \odot is the elementwise product.

2. Parameter gradients

$$\frac{\partial \mathcal{L}}{\partial W^{[l]}} = \delta^{[l]} (a^{[l-1]})^\top, \quad \frac{\partial \mathcal{L}}{\partial b^{[l]}} = \text{sum}(\delta^{[l]})$$

(sum over batch and spatial axes as appropriate).

3. Back-propagate to previous activations

$$\frac{\partial \mathcal{L}}{\partial a^{[l-1]}} = (W^{[l]})^\top \delta^{[l]}$$

(for CNNs, this is a convolution with flipped kernels).

Finally, update parameters (e.g., SGD):

$$W^{[l]} \leftarrow W^{[l]} - \eta \frac{\partial \mathcal{L}}{\partial W^{[l]}}, \quad b^{[l]} \leftarrow b^{[l]} - \eta \frac{\partial \mathcal{L}}{\partial b^{[l]}}$$

Batch setting: replace outer products/sums by their batch-averaged counterparts.

Analogy

Think of a **factory line**:

- **Forward pass:** raw material (input) flows through stations (layers) to produce a product (prediction) and a quality score (loss).
 - **Backward pass:** a **quality inspector** walks backward, station by station, telling each what adjustments reduce defects. Those adjustments are the **gradients**.
-

Mathematical Foundation

A) Chain Rule (multivariate)

For composed mappings $x \xrightarrow{f_1} u \xrightarrow{f_2} v \xrightarrow{f_3} \mathcal{L}$,

$$\frac{\partial \mathcal{L}}{\partial x} = \frac{\partial \mathcal{L}}{\partial v} \frac{\partial v}{\partial u} \frac{\partial u}{\partial x}$$

Backprop is **reverse-mode autodiff**: compute a single scalar loss gradient wrt **all** parameters in time proportional to a few forward passes.

B) Layerwise derivatives (dense layer)

For $z^{[l]} = W^{[l]}a^{[l-1]} + b^{[l]}$, $a^{[l]} = \phi^{[l]}(z^{[l]})$:

- Activation derivative:

$$\frac{\partial \mathcal{L}}{\partial z^{[l]}} = \frac{\partial \mathcal{L}}{\partial a^{[l]}} \odot \phi'^{[l]}(z^{[l]})$$

- Weights, bias:

$$\frac{\partial \mathcal{L}}{\partial W^{[l]}} = \left(\frac{\partial \mathcal{L}}{\partial z^{[l]}} \right) (a^{[l-1]})^\top, \quad \frac{\partial \mathcal{L}}{\partial b^{[l]}} = \text{sum} \left(\frac{\partial \mathcal{L}}{\partial z^{[l]}} \right)$$

- Previous activations:

$$\frac{\partial \mathcal{L}}{\partial a^{[l-1]}} = (W^{[l]})^\top \left(\frac{\partial \mathcal{L}}{\partial z^{[l]}} \right)$$

C) Common activation derivatives

ReLU: $\phi'(z) = \mathbf{1}[z > 0]$;

Sigmoid: $\sigma'(z) = \sigma(z)(1 - \sigma(z))$;

Tanh: $\tanh'(z) = 1 - \tanh^2(z)$

D) Softmax + Cross-Entropy (clean top-layer gradient)

Let $\hat{y} = \text{softmax}(z^{[L]})$, one-hot y , and $\mathcal{L} = -\sum_c y_c \log \hat{y}_c$.

Then a key simplification:

$$\delta^{[L]} \equiv \frac{\partial \mathcal{L}}{\partial z^{[L]}} = \hat{y} - y$$

This is why softmax-CE is numerically and analytically convenient.

E) Convolution layers (intuition of gradients)

- $z^{[l]} = W^{[l]} * a^{[l-1]} + b^{[l]}$ ($*$ = convolution)
- Weight gradient: **correlate** input activations with $\delta^{[l]}$.
- Input gradient: **convolve** $\delta^{[l]}$ with **flipped** kernels (transpose conv).
- Bias gradient: sum $\delta^{[l]}$ over batch and spatial positions.

F) Stability tricks

- Use **cached forward values** (e.g., $\sigma(z)$ to compute $\sigma'(z)$).
- **BatchNorm**, **residual connections**, and good inits (He/Xavier) stabilize gradients.
- **Gradient checking** (finite differences) can verify an implementation.

Use Case

Training an image classifier:

1. Forward: images \rightarrow CNN \rightarrow logits \rightarrow softmax \rightarrow cross-entropy loss.

2. Backward: compute $\hat{y} - y$ at the output, propagate through FC, conv, pool, and activations to get $\{\partial \mathcal{L}/\partial W, \partial \mathcal{L}/\partial b\}$ for **every layer**.
 3. Update params with SGD/Adam. Repeat until convergence.
-

Invariance vs. Equivariance

High-Level Summary

- A function is **invariant** to a transformation if the **output doesn't change** when the input changes in a specific way.
 - A function is **equivariant** if the **output changes in a predictable way** when the input changes.
-

Detailed Explanation

◆ 1. Invariance

- A model is **invariant** if certain transformations of the input leave the output unchanged.
- Example: A classifier should still predict "cat" whether the cat image is **shifted, rotated, or resized**.

Formally:

$$f(T(x)) = f(x)$$

Where:

- T = transformation (e.g., translation, rotation).
- f = model.

Interpretation: The model **ignores irrelevant variations**.

◆ 2. Equivariance

- A model is **equivariant** if the transformation of the input leads to a **corresponding transformation of the output**.

- Example: In semantic segmentation, if you shift the input image by 10 pixels, the **segmentation mask** should also shift by 10 pixels.

Formally:

$$f(T(x)) = T(f(x))$$

Interpretation: The model **tracks transformations consistently**.

◆ 3. Examples in Deep Learning

- **CNNs:**
 - Convolutions are **translation equivariant** (shifting input shifts feature map).
 - Pooling layers add **translation invariance** (small shifts don't change output).
 - **Classification networks:** We want **invariance** (e.g., rotated "3" is still a "3").
 - **Detection/Segmentation networks:** We want **equivariance** (shifted object → shifted bounding box/mask).
-

Analogy

- **Invariance:** Imagine a **face-recognition system** → No matter how you rotate or slightly shift your head, it still recognizes *you*.
 - **Equivariance:** Imagine **Google Maps arrows** → if the map rotates, the arrows rotate **accordingly**, preserving orientation.
-

Mathematical Foundation

- **Invariance:**

$$f(T(x)) = f(x)$$

(Output doesn't change under transformation).

- **Equivariance:**

$$f(T(x)) = T(f(x))$$

(Output transforms the same way as input).

Use Case

- **Invariance:**
 - Classification (digit recognition, face ID, object recognition).
 - **Equivariance:**
 - Segmentation, detection, pose estimation (where location/orientation matters).
-

Key Takeaway:

- **Invariance = ignore transformations.**
 - **Equivariance = respect transformations.**
-