

Transformer

Core Idea of the Transformer

Before Transformers, models like **RNNs** and **LSTMs** were the go-to for sequence tasks (translation, summarization, etc.).

However, they processed sequences **sequentially**, which caused two big problems:

1. **No parallelization** — training was slow.
2. **Difficulty capturing long-range dependencies** — information faded over long sequences.

 **Transformers** solved both using one radical idea:

 Use attention (especially self-attention) instead of recurrence to model dependencies — and process all tokens in parallel.

This allows:

- Parallel training (faster computation)
 - Global context awareness (long-range relationships)
 - Better scalability with depth and data
-



Transformer Architecture Overview

The Transformer is composed of two main parts:

Encoder → Decoder

Each is made up of **repeated blocks (layers)**.

Encoder–Decoder Overview

Component	Function
Encoder	Reads the input sentence and produces contextualized representations.
Decoder	Uses those representations to generate the output sequence (e.g., translated sentence).

For example:

| Input: "I love dogs" (English)
| Output: "J'aime les chiens" (French)

ENCODER ARCHITECTURE

Each encoder block has **two main sublayers**:

- 1 Multi-Head Self-Attention**
- 2 Feed-Forward Network (FFN)**

Plus two important add-ons:

- **Residual connection** around each sublayer
 - **Layer normalization**
-

Step-by-Step in Encoder:

1. Input Embedding

Each word/token is first mapped to a dense vector:

$$x_i \rightarrow e_i \in \mathbb{R}^{d_{model}}$$

2. Positional Encoding

Since Transformers have **no recurrence or convolution**, they don't know word order.

So we add **positional encodings** (sine & cosine patterns) to embeddings.

$$z_i = e_i + PE_i$$

This tells the model the position of each token.

3. Multi-Head Self-Attention

Every token **attends to all tokens**, learning contextual meaning.

Output = weighted sum of all token representations (as explained in self-attention).

4. Add & Norm

Output of attention is added back to the input (residual connection) → normalized:

$$\text{LayerNorm}(x + \text{Attention}(x))$$

5. Feed Forward Network (FFN)

A simple MLP applied independently to each position:

$$\text{FFN}(x) = \text{ReLU}(xW_1 + b_1)W_2 + b_2$$

6. Add & Norm again

Another residual connection + normalization.

 Final output = encoded representations of each token with full context awareness.

DECODER ARCHITECTURE

Each decoder block has **three sublayers**:

- 1 Masked Multi-Head Self-Attention**
- 2 Encoder-Decoder Attention**
- 3 Feed-Forward Network**

Again with residuals and normalization.

Step-by-Step in Decoder:

1. Masked Multi-Head Self-Attention

The decoder can only attend to **previous tokens** (to prevent "cheating" during generation).

Masking ensures attention weights for future positions = 0.

2. Encoder–Decoder Attention

Now the decoder attends to **the encoder's output** — this is how it aligns with the input sentence.

(e.g., while generating "chiens," it attends to "dogs.")

3. Feed Forward + Add & Norm

Same as encoder.

4. Linear + Softmax

Finally, the decoder outputs probabilities for the next token.

$$P(y_t | y_{<t}, X)$$

📦 The Complete Transformer Flow

Input sentence → [Encoder stack] → Context vectors

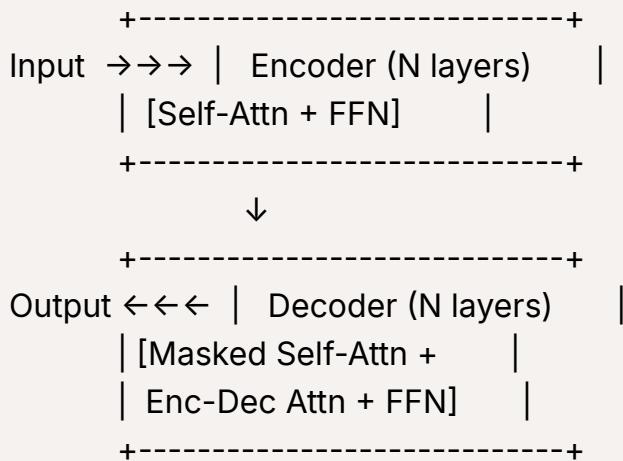
Decoder (auto-regressive) → uses context + previous outputs → generates target sequence

⚡ Key Concepts Recap

Concept	Role
Self-Attention	Lets each word see all others for context
Multi-Head Attention	Captures multiple relationships in parallel
Positional Encoding	Injects sequence order information
Feed Forward Network	Adds non-linearity and depth
Residual Connections + LayerNorm	Stabilize and speed up training
Masked Attention (Decoder)	Ensures autoregressive (left-to-right) generation

Concept	Role
Encoder–Decoder Attention	Connects source and target sequences

🧩 Visual Summary



🎯 Intuition Summary

Step	Analogy
Encoder	Reads and understands the full sentence (like a human listening carefully).
Decoder	Writes the translation step-by-step, looking back at both the source (encoder) and what it has already written.
Attention	The mechanism that decides <i>what to focus on</i> in each step.

🧠 The Goal of Transformer Training

Transformers are trained to **predict the next token** in a sequence —
That's how they learn language understanding and translation.

🎯 Objective (Training Goal)

Given:

- An **input sequence** ($X = [x_1, x_2, \dots, x_n]$)
- A **target sequence** ($Y = [y_1, y_2, \dots, y_m]$)

We train the Transformer to **predict each token** (y_t)

Given the **previous tokens** and the **input**:

$$P(y_t|y_{<t}, X)$$

The model is **auto-regressive** — it predicts tokens one by one.

Step-by-Step: Transformer Training Pipeline

1 Input Processing

1. **Input tokens** → embedded into vectors.
2. **Positional encoding** was added to keep word order.
3. Passed into the **encoder** stack (N layers).

Encoder outputs **context vectors** that summarize the input sentence.

2 Decoder Operation During Training

During training, we already know the **target sentence**.

So we feed the **ground-truth tokens** into the decoder — this is called **teacher forcing**.

Example:

If the target sentence is

| "I love pizza 🍕"

Then during training:

- Input to decoder = "I love"
- Target output = "pizza"

Masking ensures that the model only "sees" previous tokens ("I", "love")

when predicting the next one.

3 Output Prediction

The **decoder's last layer** outputs a vector of size equal to the model dimension (say 512).

This goes through a **linear layer + softmax** to produce a probability distribution over the vocabulary.

$$P(y_t|y_{<t}, X) = \text{softmax}(W_{out} h_t)$$

where (h_t) is the decoder's output at time (t).

4 Loss Function — Cross-Entropy Loss

We use **cross-entropy loss**, which compares predicted probabilities to the true token.

$$\mathcal{L} = - \sum_{t=1}^m \log P(y_t^{true}|y_{<t}, X)$$

👉 This penalizes the model more when it assigns a low probability to the correct word.

5 Backpropagation Through the Transformer

This is where training happens!

◆ Step 1: Compute Loss Gradient

The loss gradient flows **backward** from the softmax output to the final decoder layer.

◆ Step 2: Gradient Through Decoder Layers

Each decoder layer has:

1. Masked Self-Attention
2. Encoder–Decoder Attention
3. Feed-Forward Network (FFN)

Gradients pass backward through each of these sub-layers.

- The gradient updates **attention weights** (W_Q, W_K, W_V)
- Updates **FFN weights**
- And adjusts **layer norm parameters**

Residual connections ensure **smooth gradient flow** — preventing vanishing gradients even in deep stacks.

◆ Step 3: Gradient Through Encoder Layers

The gradient also flows backward into the **encoder stack**, because decoder attention depends on encoder outputs.

Each encoder layer gets updates for:

- Self-attention weights (captures better token dependencies)
- FFN weights (refines non-linear transformations)
- Layer norm and residuals

Thus, the encoder gradually learns to represent input sentences more effectively.

◆ Step 4: Gradient to Embeddings and Positional Encoding

At the end of backpropagation:

- Word embedding matrix (W_E) gets updated → better representations of words.
 - Positional encoding (if learned, not fixed sinusoidal) can also be adjusted.
-

6 Optimization

Transformers typically use the **Adam optimizer** (or AdamW) with a special **learning rate schedule**:

$$\text{lr} = d_{model}^{-0.5} \cdot \min(\text{step}^{-0.5}, \text{step} \cdot \text{warmupsteps}^{-1.5})$$

This means:

- LR increases linearly during early “warm-up” steps
- Then decays proportionally to ($\text{step}^{-0.5}$)

This helps stabilize training in the beginning and improve convergence later.

7 Training Objective in Practice

The model's training objective over the full dataset:

$$\text{Minimize } \mathcal{L}_{total} = - \sum_{(X, Y) \in D} \sum_{t=1}^{|Y|} \log P(y_t | y_{<t}, X; \theta)$$

where (θ) includes **all learnable parameters**:

- Embedding weights
 - Attention weights ((W_Q, W_K, W_V, W_O))
 - FFN weights ((W_1, W_2))
 - Layer norms
 - Output projection (W_{out})
-

How Each Layer Learns (Conceptually)

Layer	What It Learns
Embedding	Word meaning (distributed representation)
Positional Encoding	Word order information
Encoder Self-Attention	Relationships among input tokens
Decoder Self-Attention	Relationships among generated tokens
Encoder–Decoder Attention	Alignments between input and output (like translation pairs)
Feed-Forward Networks	Complex nonlinear transformations of contextual info
Output Softmax Layer	Vocabulary-level probability mapping

All layers are trained **jointly** — the loss from the final output is backpropagated to all components.

Training Example (Simplified)

Suppose your training pair is:

| Input: "I love dogs"

| Output: "J'aime les chiens"

At step **t=3**:

- Model predicts token = "chiens"
- True token = "chiens"
- Cross-entropy loss is low → minimal update

If model predicted "chat" (cat):

- Loss is high → gradient adjusts encoder-decoder attention so "dogs" better maps to "chiens"

Over many samples, these gradients train:

- Encoders to encode semantic meaning
- Decoders to decode contextually correct translations

🎯 Intuitive Summary

Step	What Happens
Forward pass	Model predicts next token using all attention layers
Compute loss	Compare predicted vs. true token (cross-entropy)
Backward pass	Gradients flow through decoder → encoder → embeddings
Update weights	Optimizer adjusts parameters
Repeat	Until the model converges and can generate accurate sequences

🧠 Big Picture: The Transformer Has Two Main Parts

[Encoder Stack] → [Decoder Stack]

But depending on the **task**, we can use:

- only the **Encoder** part,

- only the **Decoder** part,
 - or both together (**Encoder–Decoder**).
-

1 Encoder-Only Models

Architecture

Use **only the encoder stack** from the Transformer.

Each layer contains:

- Multi-head **self-attention**
- Feed-forward network
- Add & Norm connections

How It Works

- The encoder takes an input sequence (like a sentence or document).
- Each token attends to *all other tokens* (bidirectionally).
- The model learns **contextual representations** of the entire input.

$$h_i = f(x_1, x_2, \dots, x_n)$$

So each token's vector (h_i) knows the meaning of all words around it.

Use Cases

Encoder-only models are used for **understanding tasks** (not generation).

Examples:

- Sentence classification (e.g., sentiment analysis)
- Named Entity Recognition (NER)
- Question answering (extractive)
- Similarity and embedding generation

Examples of Encoder-Only Models

Model	Description
BERT	"Bidirectional Encoder Representations from Transformers" — learns deep bidirectional context.
RoBERTa	Robustly optimized version of BERT.
DistilBERT	Smaller, faster version of BERT.

⚡ Key Property

Bidirectional attention: each token can see all other tokens on both sides — left and right.

This gives rich contextual understanding but makes **text generation impossible** (since the model "sees the future").

2 Decoder-Only Models

🧩 Architecture

Use **only the decoder stack**, but **without encoder-decoder attention**.

Each layer includes:

- **Masked self-attention**
- **Feed-forward network**
- Add & Norm connections

🔍 How It Works

- The decoder predicts the next token one step at a time.
- Masked attention ensures each token only attends to **previous tokens** (not future ones).

$$P(y_t | y_{<t})$$

This creates a **causal**, left-to-right generation process.

🧠 Use Cases

Decoder-only models are used for **generation tasks** such as:

- Text completion
- Dialogue and chatbots
- Story generation
- Code generation
- Autoregressive modeling

Examples of Decoder-Only Models

Model	Description
GPT (1, 2, 3, 4, 5)	"Generative Pre-trained Transformer" — trained to predict next word (language modeling).
LLaMA, Falcon, Mistral	Open-source GPT-style models.
CodeGen, StarCoder	Specialized for code generation.

Key Property

Unidirectional attention: each token can only see tokens **to its left**, preserving causality.

This makes it perfect for **autoregressive generation**.

3 Encoder–Decoder (Seq2Seq) Models

Architecture

Uses **both encoder and decoder stacks** — the *full* Transformer.

Encoder → produces context → Decoder → generates output

How It Works

- The **encoder** processes the input sequence → context representations.
- The **decoder** uses:
 - Masked self-attention (to generate outputs step-by-step)

- Encoder–decoder attention (to focus on relevant input tokens)
- This allows **conditional generation** (output depends on input).

$$P(y_t|y_{<t}, X)$$

Use Cases

Used for **sequence-to-sequence tasks**, where input and output are different sequences:

- Machine translation
- Summarization
- Text-to-text transformation
- Question answering (generative)
- Paraphrasing

Examples of Encoder–Decoder Models

Model	Description
T5	“Text-To-Text Transfer Transformer” — converts all NLP tasks into a text-to-text format.
BART	Combines BERT-style encoder + GPT-style decoder for text generation and denoising.
MarianMT	Specialized for machine translation.
mT5, Flan-T5	Multilingual or instruction-tuned versions.

Key Property

Bidirectional in the encoder, unidirectional in the decoder.

→ Model *understands* input deeply, and *generates* conditioned on it.

Comparison Summary Table

Feature	Encoder-Only	Decoder-Only	Encoder-Decoder
Attention Direction	Bidirectional	Unidirectional (causal)	Encoder: bidirectional Decoder: unidirectional
Main Purpose	Understanding	Generation	Translation / Seq2Seq
Inputs	Single text	Previous tokens	Input + Generated output
Examples	BERT, RoBERTa	GPT, LLaMA	T5, BART
Use Cases	Classification, QA (extractive)	Text completion, chatbots	Summarization, translation
Training Objective	Masked LM (fill missing words)	Next-token prediction	Conditional generation
Context Flow	All tokens see each other	Each token sees past only	Decoder attends to encoder outputs

Visual Summary

1 Encoder-Only

Input → [Encoder Stack] → Output Representation

↳ Understanding task (BERT)

2 Decoder-Only

Input → [Masked Decoder Stack] → Generated Output

↳ Generation task (GPT)

3 Encoder-Decoder

Input → [Encoder] → Context → [Decoder] → Output

↳ Translation / Summarization (T5, BART)

Intuition Summary

Type	Analogy
Encoder-Only	Like reading and <i>understanding</i> a sentence deeply.

Type	Analogy
Decoder-Only	Like <i>writing</i> a story word-by-word.
Encoder-Decoder	Like <i>translating</i> — read a source sentence, then generate its version in another language.



What Is Pretraining in Transformers?

Before a Transformer can perform tasks like translation, summarization, or sentiment analysis, it needs to *understand* language.

To gain this understanding, it undergoes a **pretraining phase** — learning from massive amounts of unlabeled text using **self-supervised objectives** (like predicting missing words).

These objectives are called **Pretraining Strategies**.



Main Pretraining Strategies in Transformers

Below are the most common and influential pretraining strategies used in Transformer models:

1. Masked Language Modeling (MLM) — Used by BERT

Idea:

- Randomly mask (hide) some words in a sentence.
- Ask the model to predict those masked words from the surrounding context.

Example:

Input: "The cat sat on the [MASK]."

Target: "mat"

Goal:

Learn *bidirectional context* — i.e., understand words based on *both left and right* neighbors.

Used in:

2. Next Sentence Prediction (NSP)

Idea:

- Alongside MLM, the model also learns whether two sentences logically follow each other.

Example:

Sentence A: "The cat sat on the mat."

Sentence B: "It started to purr."  (Next sentence)

Sentence C: "Apples grow on trees."  (Not the next sentence)

Goal:

Learn *relationships between sentences* — helpful for question answering and natural language inference.

Used in:

 Original BERT

Limitation:

Later research (e.g., RoBERTa) showed NSP doesn't help much and can be removed.

3. Causal Language Modeling (CLM) — Used by GPT

Idea:

- Predict the *next word* given all previous words.
- Only uses *left-to-right* context (unidirectional).

Example:

Input: "The cat sat on the"

Target: "mat"

Goal:

Learn *generative* modeling — essential for text generation and completion.

Used in:

 GPT, GPT-2, GPT-3, GPT-4, LLaMA

4. Permutation Language Modeling (PLM) — Used by XLNet

Idea:

- Instead of masking words, predict tokens in a *random permutation order*.
- This combines the benefits of MLM and CLM (bidirectional context + generative ability).

Used in:

 XLNet

Goal:

Capture bidirectional context *without using masks*.

5. Denoising Autoencoder (DAE) — Used by BART / T5

Idea:

- Corrupt the input sentence (by masking, deleting, shuffling words).
- Ask the model to *reconstruct* the original sentence.

Example:

Corrupted: "The [MASK] on mat cat the."

Target: "The cat sat on the mat."

Goal:

Learn to recover meaning from noisy input — great for summarization, translation, etc.

Used in:

 BART, T5

Summary Table

Strategy	Directionality	Objective	Example Models	Strength
MLM	Bidirectional	Predict masked words	BERT	Strong contextual understanding
NSP	Bidirectional	Predict next sentence	BERT	Sentence-level reasoning
CLM	Unidirectional	Predict next token	GPT series	Natural text generation
PLM	Bidirectional (permuted)	Predict token order	XLNet	Combines BERT & GPT benefits
DAE	Bidirectional	Reconstruct corrupted input	BART, T5	Robust understanding & generation

Effect on Training and Downstream Tasks

- These strategies help the model **learn general language representations** from unlabeled data.
- During **fine-tuning**, the pretrained weights are adjusted slightly for specific tasks like:
 - Classification (sentiment)
 - QA
 - Summarization
 - Translation
- This approach drastically reduces labeled data requirements and training time.

What is Masked Language Modeling (MLM)?

Definition:

Masked Language Modeling is a **self-supervised learning objective** where a model learns to **predict missing (masked) words** in a sentence based on their surrounding context.

Goal

Instead of predicting the next word (like GPT), MLM teaches the model to **understand context in both directions** — left and right.

That's why it's called a **bidirectional training objective**.

Example

Original sentence:

| "The cat sat on the mat."

We randomly **mask** one or more tokens (e.g., 15% of them):

| "The cat sat on the [MASK]."

The model must predict the missing word:

| "mat"

So it learns to understand *how words relate to each other in both directions*.

How It Works (Step-by-Step)

1 Input Preparation

- Take a sentence and **randomly mask** 15% of the tokens.
- But not all of them are replaced with **[MASK]**:
 - 80% → replaced with **[MASK]**
 - 10% → replaced with a random word
 - 10% → left unchanged

This helps prevent the model from overfitting to the **[MASK]** token.

Example:

Original	Masked Input	Target
"I love NLP models."	"I love [MASK] models."	"NLP"

2 Encoder Processing

The masked sentence is passed through the **encoder** (e.g., in BERT):

- Every token attends to *all* other tokens (including left and right context).
- The encoder produces contextual embeddings for each token.

3 Prediction Layer

For each masked position, the model predicts the **original word** using a softmax classifier over the vocabulary:

$$P(w_i | \text{context}) = \text{softmax}(W h_i + b)$$

where:

- (h_i) = hidden representation of the masked position
- (W) = output projection matrix
- (b) = bias vector

4 Loss Function — Cross-Entropy Loss

The model is trained to minimize the negative log-likelihood of the correct token:

$$\mathcal{L} = - \sum_{i \in M} \log P(w_i^{\text{true}} | \text{context})$$

where (M) = set of masked positions.

Only masked tokens contribute to the loss.

Intuitive Understanding

Property	Description
Bidirectional context	Model looks at both left and right sides of the masked word.
Self-supervised	Labels are created from the data itself (no manual annotation needed).
Contextual embeddings	Learns meaning of words <i>in context</i> (e.g., "bank" in "river bank" vs "money bank").

🏗 Example in Detail

Sentence:

| "The dog chased the [MASK]."

The model sees:

- Left context: "The dog chased the"
- Right context: (none in this case)

Predicts:

| "ball" (high probability), "cat", "stick" (lower probability)

⚡ Why MLM Works So Well

- It teaches the model to **understand relationships among all words** in a sentence.
- It's **bidirectional** — unlike autoregressive models (like GPT) which only look left-to-right.
- The representations learned can be reused for **many downstream tasks** (transfer learning).

📘 Example: BERT's MLM Training Objective

BERT combines two tasks during pretraining:

Task	Description
Masked Language Modeling (MLM)	Predict masked tokens using bidirectional context.
Next Sentence Prediction (NSP)	Predict if one sentence follows another.

During training, BERT randomly masks 15% of input tokens and learns to predict them.

After pretraining, it's fine-tuned for specific NLP tasks (classification, QA, etc.).

Comparison to Next Token Prediction (like GPT)

Feature	Masked LM (BERT)	Next-Token LM (GPT)
Context	Bidirectional	Left-to-right (causal)
Masking	Predict missing words	Predict next word
Use Case	Understanding	Generation
Example Task	Fill-in-the-blank	Text continuation

Summary

Aspect	Masked Language Modeling
What it does	Randomly masks words and trains model to predict them
Why it works	Forces model to learn bidirectional contextual understanding
Loss function	Cross-entropy over masked tokens
Used in	BERT, RoBERTa, ALBERT
Result	Powerful contextual embeddings for downstream NLP tasks