

# Linear Regression I

---

## Linear Regression Hypothesis

---

### 1. Detailed Conceptual Overview

Linear Regression is one of the most fundamental algorithms in Machine Learning. It assumes that the relationship between the input features (**independent variables**) and the output (**dependent variable**) can be approximated using a **linear function**.

**Core Idea:**

- Predict  $y$  as a **linear combination of input features  $x$** .
- Example: Predict house price using features like size, age, and location score.

**General Hypothesis Form:**

- For one feature (simple regression):

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Here,  $\theta_0$  = intercept (bias),  $\theta_1$  = slope (weight).

- For multiple features (multivariate regression):

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d$$

Compact form (vectorized):

$$h_{\theta}(x) = \theta^T x$$

where  $\theta$  = parameter vector,  $x$  = feature vector (with a "1" appended for bias).

**Interpretation:**

- Each coefficient  $\theta_j$  represents how much the target  $y$  changes when feature  $x_j$  increases by 1, while keeping other features constant.
- The intercept  $\theta_0$  represents the baseline value of  $y$  when all features are zero.

## Examples:

1. **Simple Regression (1D):** Predicting salary based on years of experience.
  - Line:  $\text{Salary} = \text{Base} + (\text{Rate} \times \text{Years})$ .
2. **Multiple Regression (2D+):** Predicting house price from size and number of rooms.
  - Plane in 3D space:  $\text{Price} = \text{Intercept} + w_1 \cdot \text{Size} + w_2 \cdot \text{Rooms}$ .

The linear regression hypothesis forms the **starting point for many ML models** (generalized linear models, logistic regression, neural networks, all extend from here).

---

## 2. Key Takeaways

- Linear Regression assumes a **linear relationship** between inputs and outputs.
  - Hypothesis is a **weighted sum of input features plus an intercept**.
  - **Simple Regression:** One input → a line.
  - **Multiple Regression:** Many inputs → a hyperplane in higher dimensions.
  - Coefficients ( $\theta$ ) capture the **influence of each feature**.
  - Baseline intercept ( $\theta_0$ ) shifts the line/plane vertically.
  - Forms the **foundation for more advanced ML models**.
- 

# Linear Regression Illustration (1D, 2D, Higher Dimensions)

---

## 1. Detailed Conceptual Overview

To better understand **linear regression**, it helps to **visualize how the hypothesis looks in different dimensions**:

---


### a) 1D Case (Simple Linear Regression):

- One input variable  $x$ .
- Hypothesis:  

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$
- Graph: A **straight line** in 2D space ( $x$ -axis = input,  $y$ -axis = output).
- Example: Predicting salary based on years of experience.

👉 Interpretation: Increasing  $x$  by 1 increases prediction by slope ( $\theta_1$ ).

$$\mathbf{x} = (x) \in \mathbb{R}$$



Vectors:  
 $w = [w_0, w_1, \dots, w_d]$   
 $x = [1, x_1, \dots, x_d]$


### b) 2D Case (Two Input Features):

- Two input variables  $(x_1, x_2)$ .
- Hypothesis:  

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$
- Graph: A **plane** in 3D space.
- Example: Predicting house price based on size ( $x_1$ ) and number of rooms ( $x_2$ ).

👉 Interpretation: Each coefficient adjusts the slope of the plane along that feature axis.

$$\mathbf{x} = (x_1, x_2) \in \mathbb{R}^2$$



### c) Higher Dimensions (Multivariate Regression):

- Many features  $(x_1, x_2, \dots, x_d)$ .
- Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

- Graph: A **hyperplane** in d+1 dimensional space.
  - Cannot visualize beyond 3D, but mathematically, the same principle applies.
  - Example: Predicting medical costs using features like age, BMI, smoking habits, and number of dependents.
- 

#### Role of the Intercept ( $\theta_0$ ):

- Shifts the line/plane/hyperplane up or down.
  - Ensures predictions aren't forced to pass through the origin (0,0,...,0).
  - Example: Even if a house has size = 0, the intercept may represent baseline cost (land value, taxes, etc.).
- 

## 2. Key Takeaways

- **1D Linear Regression:** Line fit between input & output.
  - **2D Linear Regression:** Plane fit between two inputs & output.
  - **Higher-Dimensional Regression:** Hyperplane fit for many features.
  - Visualization becomes harder as dimensions increase, but math is the same.
  - Intercept shifts the regression surface vertically.
  - Each feature contributes a **slope** (weight), controlling how strongly it affects predictions.
- 

# Choosing the Best Coefficients (Error Minimization in Linear Regression)

---

## 1. Detailed Conceptual Overview

In Linear Regression, our hypothesis is:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

But how do we **find the best coefficients ( $\theta$ )?**

The idea:

- The model should make predictions **as close as possible to actual outputs**.
  - We measure closeness using an **error function (loss function)**.
  - The "best" coefficients are those that **minimize this error across the training dataset**.
- 

**How it works:**

1. Take all training points  $(x_i, y_i)$ .

2. Compute prediction error:

$$\text{error}_i = y_i - h_\theta(x_i)$$

3. Aggregate errors into a **cost function** (commonly **Sum of Squared Errors**):

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (y_i - h_\theta(x_i))^2$$

4. Find  $\theta$  that minimizes  $J(\theta)$ .

👉 This is essentially an **optimization problem**.

- For small problems → we can solve analytically (normal equation).
  - For large problems → we use iterative optimization (like Gradient Descent).
- 

**Why Squared Error?**

- Penalizes larger mistakes more heavily.
  - Makes the function smooth and differentiable, useful for optimization.
  - Leads to a unique solution if features are independent.
- 

## 2. Key Takeaways

- Best coefficients are chosen by **minimizing prediction error**.
- Error is usually measured as **sum of squared differences** between predicted and actual values.
- This leads to a **cost function  $J(\theta)$** , which is minimized to estimate  $\theta$ .
- Two approaches to minimization:

- **Analytical solution** (Normal Equation).
  - **Iterative solution** (Gradient Descent).
  - Error minimization ensures that the chosen line/plane/hyperplane is the "**best fit**" for the data.
- 

## Error Metrics in Regression (L1, L2, Max Error)

---

### 1. Detailed Conceptual Overview

When fitting a regression model, we need a way to **measure how good the predictions are** compared to actual values. Different **error metrics** capture different aspects of model performance.

---

#### a) L1 Norm (Absolute Error):

- Definition: Sum of absolute differences between predictions and actuals.

$$L1 = \sum_{i=1}^n |y_i - h_\theta(x_i)|$$

- Properties:
    - More **robust to outliers** (does not exaggerate large errors).
    - Leads to the **Median** as the best-fit estimator in statistics.
  - Example usage: Lasso regression (adds L1 penalty).
- 

#### b) L2 Norm (Squared Error):

- Definition: Sum of squared differences between predictions and actuals.

$$L2 = \sum_{i=1}^n (y_i - h_\theta(x_i))^2$$

- Properties:
  - Punishes **large errors more strongly** (since errors are squared).
  - Leads to the **Mean** as the best-fit estimator.
  - Most common choice for regression (basis of least squares).

- Example usage: Ridge regression (adds L2 penalty).
- 

### c) Max Error (Infinity Norm):

- Definition: Largest absolute error across all predictions.

$$L_{\infty} = \max_i |y_i - h_{\theta}(x_i)|$$

- Properties:
    - Focuses on the **worst-case error**.
    - Useful when even a single large mistake is unacceptable (e.g., safety-critical systems).
  - Example usage: Quality control, robotics.
- 

### Comparison:

- L1: Robust to outliers, but may give less stable solutions.
  - L2: Sensitive to outliers, but gives smooth, unique solutions.
  - Max Error: Ensures no extreme failure, but ignores average performance.
- 

## 2. Key Takeaways

- **Error metrics quantify model performance in regression.**
  - **L1 norm** = absolute error → robust to outliers.
  - **L2 norm** = squared error → emphasizes large errors, most common.
  - **Max error** = worst-case error → ensures safety-critical reliability.
  - Choice of metric depends on application (robustness vs sensitivity vs safety).
- 

# Sum of Squared Errors (SSE) & In-sample vs Out-of-sample Error

---

## 1. Detailed Conceptual Overview

When fitting a regression model, we need a way to evaluate how well the model explains the data. Two key ideas here are **Sum of Squared Errors (SSE)** and the distinction between **In-sample vs Out-of-sample error**.

---

### a) Sum of Squared Errors (SSE):

- Definition: Total squared difference between predicted and actual outputs across all training data.

$$SSE = \sum_{i=1}^n (y_i - h_\theta(x_i))^2$$

- Intuition:
    - Measures the **total error** the model makes on the training set.
    - Squaring emphasizes large deviations more strongly.
  - Goal: Minimize SSE when choosing coefficients ( $\theta$ ).
  - Example: If predictions match actuals perfectly  $\rightarrow$  SSE = 0.
- 

### b) In-sample Error (Training Error):

- The error the model makes on the **same data it was trained on**.
- Usually **lower**, since the model has already seen this data.
- Risk: Can be misleading if the model overfits (memorizes training data).
- Formula:

$$E_{in} = \frac{1}{n_{train}} \sum_{i=1}^{n_{train}} (y_i - h(x_i))^2$$

---

### c) Out-of-sample Error (Test/Generalization Error):

- The error the model makes on **new, unseen data**.
- More realistic measure of how well the model will perform in practice.
- Usually **higher than in-sample error**, since the model hasn't seen the test data.

- Formula:

$$E_{out} = \mathbb{E}_{(x,y) \sim P} [(y - h(x))^2]$$

---

**Key Idea:**

- A good model should have **both low in-sample and low out-of-sample error**.
  - If in-sample error is very low but out-of-sample error is very high → **overfitting**.
  - If both errors are high → **underfitting**.
- 

## 2. Key Takeaways

- **SSE (Sum of Squared Errors)**: Main cost function used in regression. Squaring emphasizes large deviations.
  - **In-sample error**: Model error on training data. It can be artificially low due to overfitting.
  - **Out-of-sample error**: Model error on unseen data. True test of generalization.
  - ML models must balance between the two: low training error **and** low test error.
  - Gap between in-sample and out-of-sample error is a strong indicator of **overfitting**.
- 

# Methodology of Regression Solutions (Fitting with 1, 2, or More Points – Least Squares Method)

---

## 1. Detailed Conceptual Overview

Linear Regression is about finding the **best-fitting line (or hyperplane)** that explains the relationship between input(s) and output. The methodology depends on how many data points we have.

---

### a) Fitting with 1 Point:

- If we only have **one data point**, infinitely many lines can pass through it.
  - The problem is **underdetermined** → we cannot uniquely define a regression line.
  - Example: Point (2, 5). Lines like  $y = 2x + 1$  or  $y = -x + 7$  both fit perfectly.
- 

### b) Fitting with 2 Points:

- If we have exactly **two data points**, there is exactly **one unique straight line** that passes through both.
  - Regression is trivial here → line is perfectly determined.
  - Example: Points (1, 2) and (3, 6) → slope =  $(6-2)/(3-1) = 2$ , so line =  $y = 2x$ .
- 

### c) Fitting with More than 2 Points:

- With more than two points, in general, **no single line will pass through all points**.
  - We must find the **best-fitting line** that minimizes errors.
  - Solution: **Least Squares Method**.
- 

#### Least Squares Method (Core Idea):

- Goal: Minimize total squared error between predictions and actual values.

$$J(\theta) = \sum_{i=1}^n (y_i - h_\theta(x_i))^2$$

- Why squares?
    - Eliminates negatives (otherwise positive and negative errors cancel).
    - Amplifies large errors → line fits the bulk of data, not just extremes.
  - Result: A unique line that minimizes SSE.
- 

#### Geometric Interpretation:

- Each data point has a vertical distance (error) from the regression line.

- Least Squares finds the line that minimizes the sum of **squared vertical distances**.
- 

## 2. Key Takeaways

- **1 Point:** Cannot define regression line (underdetermined).
  - **2 Points:** Unique line fits perfectly.
  - **>2 Points:** Use the **Least Squares Method** to find the best fit.
  - **Least Squares Principle:** Minimize total squared error between predictions and actuals.
  - This methodology generalizes to multiple regression (planes/hyperplanes).
  - Core idea: Regression is about finding the line/hyperplane that **balances errors across all points**.
- 

# Algorithmic Solution (Matrix Pseudo-inverse Method)

---

## 1. Detailed Conceptual Overview

Up to now, we've seen regression as finding the **best-fit line/hyperplane** by minimizing the **sum of squared errors (SSE)**. But how do we **actually compute the coefficients ( $\theta$ )?**

One powerful approach is the **Matrix Pseudo-inverse Method** — also known as the **Normal Equation Solution**.

---

### a) Setup:

- Data matrix:

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & x_{1d} \\ 1 & x_{21} & x_{22} & \dots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & x_{n2} & \dots & x_{nd} \end{bmatrix}$$

where each row = one data point, first column = 1 (for intercept).

- Target vector:

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

- Model prediction:

$$\hat{y} = X\theta$$

---

#### b) Least Squares Objective:

- We want to minimize:

$$J(\theta) = \|y - X\theta\|^2$$

- Setting derivative = 0 leads to **Normal Equations**:

$$X^T X\theta = X^T y$$

---

#### c) Solution via Pseudo-inverse:

- If  $X^T X$  is invertible:

$$\theta = (X^T X)^{-1} X^T y$$

- If  $X^T X$  is **not invertible** (e.g., collinear features, more features than data points), we use the **Moore-Penrose Pseudo-inverse**:

$$\theta = X^+ y$$

where  $X^+$  = pseudo-inverse of X.

---

#### d) Practical Aspects:

- Works well for **small to medium datasets** (analytical closed-form solution).
- Computational cost:  $O(d^3)$  for matrix inversion.
- For **large datasets**, iterative methods (like gradient descent) are more efficient.

## 2. Key Takeaways

- Regression coefficients can be solved directly using **matrix algebra**.
  - Core formula:
$$\theta = (X^T X)^{-1} X^T y$$
  - If  $X^T X$  is not invertible → use **pseudo-inverse**.
  - Advantage: Exact solution, no need for tuning.
  - Limitation: Expensive for high-dimensional data → iterative methods preferred.
  - Foundation for many ML algorithms — even advanced ones often rely on pseudo-inverse under the hood.
- 

# Generalization of Regression Models

---

## 1. Detailed Conceptual Overview

So far, we have focused on **fitting a regression to training data**. But in ML, the true test is **how well the model generalizes** — i.e., performs on unseen data.

**Generalization** = the ability of a regression model to make accurate predictions on **new data not used during training**.

---

### a) In-sample vs Out-of-sample Recall

- **In-sample error (training error)**: Error on the training set.
  - **Out-of-sample error (test/generalization error)**: Error on new, unseen data.
  - A model may have **very low in-sample error** but **very high out-of-sample error** → classic sign of **overfitting**.
- 

### b) Bias-Variance Tradeoff (conceptual here):

- **High bias (underfitting)**: Model too simple → fails to capture trends (both training and test errors high).
- **High variance (overfitting)**: Model too complex → memorizes training data but fails on test data (train error low, test error high).

- Good generalization requires a **balance**.
- 

### c) Factors Affecting Generalization in Regression:

#### 1. Model Complexity:

- Linear regression assumes a linear relationship.
- If the true relationship is highly nonlinear, linear regression will underfit.

#### 2. Number of Features vs Samples:

- Too many features with too few samples → overfitting.
- Dimensionality reduction or regularization helps.

#### 3. Quality of Data:

- Noisy data hurts generalization.
- Outliers can distort the regression fit.

#### 4. Validation Strategy:

- Cross-validation helps estimate generalization error.
  - Splitting into train/validation/test sets is critical.
- 

### d) Improving Generalization:

- **Regularization** (penalize large coefficients: Ridge, Lasso).
  - **Feature selection** (remove irrelevant features).
  - **Data augmentation** or collecting more training samples.
  - **Cross-validation** to tune complexity before final testing.
- 

## 2. Key Takeaways

- **Generalization = the ability to perform well on unseen data.**
- **Overfitting:** Low training error, high test error.
- **Underfitting:** Both training and test errors are high.
- Achieving good generalization requires balancing **bias and variance**.

- Strategies: regularization, feature selection, cross-validation, and more data.
  - Generalization is the **ultimate goal** of regression and all ML models.
- 

# Binary Classification with Linear Regression (0/1 Error vs Squared Error)

---

## 1. Detailed Conceptual Overview

Linear regression was designed for **predicting continuous values**. But can we use it for **classification** (e.g., spam vs not spam, churn vs no churn)?

Let's explore:

---

### a) Idea of Using Linear Regression for Classification

- Suppose labels are binary:  $y \in \{0, 1\}$ .
- We can try to fit a regression model:  
$$h_{\theta}(x) = \theta^T x$$
- Then classify based on a threshold:
  - If  $h_{\theta}(x) \geq 0.5$ , predict class **1**.
  - Otherwise, predict class **0**.

This works in some simple cases, but has major issues.

---

### b) Error Metrics for Classification

#### 1. 0/1 Error (Classification Error):

- Definition: 0 if predicted class = true class, 1 otherwise.
- Total error = count of misclassified points.
- Problem: This function is **non-differentiable** → can't be optimized with calculus.

$$\text{Error}_{0/1} = \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{h_\theta(x_i) \neq y_i\}$$

### 1. Squared Error (Regression-style Error):

- Definition: Treat  $y$  as 0/1, and minimize squared error like regression.
- Cost function:

$$J(\theta) = \sum_{i=1}^n (y_i - h_\theta(x_i))^2$$

- This is differentiable and solvable.

#### Problem:

- Predictions are not restricted to  $[0, 1]$ .
- The fitted line may produce values less than 0 or greater than 1, which don't make sense as probabilities.
- Leads to poor classification boundaries in practice.

### c) Why We Don't Use Linear Regression for Classification

- Predictions can fall outside  $[0, 1]$ .
- The decision boundary may not separate classes well.
- 0/1 error is the correct measure for classification, but it's not optimizable directly.
- Squared error is optimizable, but not ideal for classification.

👉 This is exactly why **Logistic Regression** was developed — it uses the **sigmoid function** to map predictions into  $[0, 1]$ , and optimizes using **cross-entropy loss**, which works better for classification.

## 2. Key Takeaways

- Linear regression can be **forced into classification**, but it's not a good fit.
- **0/1 error** is the natural metric for classification, but not differentiable.

- **Squared error** is optimizable, but doesn't handle classification properly.
  - Linear regression may give outputs outside  $[0, 1]$ , making interpretation problematic.
  - This motivates **Logistic Regression**, which fixes these issues by using the sigmoid and cross-entropy loss.
-