

# Linear Regression II: Generalization and Statistical Inference

---

## Methodology Recap: Fitting Regression with 1, 2, $\geq 3$ Points in Higher Dimensions

---

### 1. Detailed Conceptual Overview

This is a continuation of what we saw in **Linear Regression I** — but now extended into **higher-dimensional spaces**.

The methodology for regression depends on the **number of data points relative to the number of features (dimensions)**.

---

#### a) Case 1: 1 Point

- With only one data point, infinitely many regression lines/hyperplanes can fit.
  - The system is **underdetermined** — no unique solution.
  - Example: In 2D, one point doesn't define a line.
- 

#### b) Case 2: 2 Points

- Two points in 2D define a **unique straight line**.
  - In higher dimensions, two points define a line (but not a plane/hyperplane).
  - Still underdetermined if we want to fit a full regression with multiple features.
- 

#### c) Case 3: $\geq 3$ Points

- When we have at least as many points as required, we can fit regression meaningfully.

- But typically, with  $\geq 3$  points, **no perfect fit exists** (points don't all fall on one line/plane).
  - Regression solves this by finding the **best-fit hyperplane** that minimizes error.
- 

#### d) Higher Dimensions (Multivariate Regression)

- Inputs:  $x \in \mathbb{R}^d$ .
- Outputs:  $y \in \mathbb{R}$ .
- Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_d x_d$$

- For  $n$  data points and  $d$  features:
    - If  $n < d \rightarrow$  underdetermined (many possible solutions).
    - If  $n = d \rightarrow$  exact solution may exist (if data is well-conditioned).
    - If  $n > d \rightarrow$  regression finds a **best-fit solution** using least squares.
- 

#### e) General Rule of Thumb:

- You need **more data points than features** for stable regression.
  - Otherwise, the solution may be non-unique or unstable (high variance in coefficients).
- 

## 2. Key Takeaways

- **1 point:** No unique regression (underdetermined).
- **2 points (2D):** Unique line, but not generalizable.
- **$\geq 3$  points:** Typically, no perfect fit  $\rightarrow$  regression finds best-fit hyperplane.
- In higher dimensions:
  - If  $n < d$ : underdetermined (infinite solutions).
  - If  $n = d$ : unique solution possible.
  - If  $n > d$ : least squares finds the best-fit.

- **Golden rule:** More data points than features are needed for meaningful regression.
- 

# Generalization in Machine Learning (In-sample vs Out-of-sample Error, Higher Dimensions)

---

## 1. Detailed Conceptual Overview

The ultimate goal of regression (and ML in general) is not just to fit the training data, but to **generalize** — i.e., perform well on unseen data.

This section extends the concept of **generalization** we saw earlier, but now with emphasis on **higher-dimensional settings**.

---

### a) In-sample Error (Training Error):

- Error computed on the **training data**.
- Regression aims to minimize this via least squares.
- Typically **low** because the model has already seen this data.
- Formula:

$$E_{in} = \frac{1}{n_{train}} \sum_{i=1}^{n_{train}} (y_i - h(x_i))^2$$

---

### b) Out-of-sample Error (Test/Generalization Error):

- Error on **new, unseen data** from the same distribution.
- True measure of model quality.
- Formula (expected error):

$$E_{out} = \mathbb{E}_{(x,y) \sim P} [(y - h(x))^2]$$

---

### c) Generalization Challenge in Higher Dimensions:

- As the **number of features ( $d$ )** increases, the model becomes more flexible.
  - This often reduces in-sample error (model fits training data better).
  - But risk of **overfitting** increases → out-of-sample error may rise.
- 

### d) Geometric View:

- In low dimensions, we project the target vector  $\mathbf{y}$  onto a line/plane.
  - In high dimensions, we project onto a subspace spanned by many features.
  - If the subspace is too large relative to data size, projection becomes unstable (sensitive to noise).
- 

### e) Balancing Generalization:

- Need to avoid both **underfitting** (high bias) and **overfitting** (high variance).
  - More features ≠ always better.
  - Rule: We need **sufficiently more data points than features** for the regression to generalize well.
- 

## 2. Key Takeaways

- **In-sample error:** Training error, usually low.
  - **Out-of-sample error:** Test error, true measure of generalization.
  - In higher dimensions:
    - More features reduce training error but risk overfitting.
    - Out-of-sample error increases if the model memorizes the training data.
  - **Projection perspective:** Regression projects data onto subspaces; too many features make the projection unstable.
  - **Goal:** Balance bias (underfitting) and variance (overfitting) for good generalization.
-

# Using Error Metrics to Generalize (Matrix Solution, Noise in Data, Probability Bounds)

---

## 1. Detailed Conceptual Overview

So far, we've described the **in-sample vs out-of-sample error** conceptually.

This topic digs deeper into **how regression error behaves mathematically**, especially when noise and higher dimensions come into play.

---

### a) Matrix Solution for Regression Error

- Recall regression solution:

$$\hat{\theta} = (X^T X)^{-1} X^T y$$

- Prediction:

$$\hat{y} = X\hat{\theta}$$

- Residual error vector:

$$e = y - \hat{y}$$

- In least squares, this error is **orthogonal** to the column space of X.
- 

### b) Noise in Data

Assume data comes from:

$$y = X\theta^* + \epsilon$$

where  $\theta^*$  is the true parameter vector and  $\epsilon$  is random noise.

---

#### 1. y (Target Vector):

- The actual outputs/labels we observe.
  - Example: House prices, student scores, etc.
  - Size:  $n \times 1$  (where  $n$  = number of data points).
-

## 2. X (Design Matrix):

- Contains all input features.
- Each row = one data point, each column = one feature (plus a column of 1's for intercept).
- Size:  $n \times d$  ( $n$  = data points,  $d$  = features).

Example with 2 features:

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} \\ 1 & x_{21} & x_{22} \\ \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} \end{bmatrix}$$

---

## 3. $\theta^*$ (True Coefficients):

- The real, underlying weights that nature uses to generate  $y$ .
- Regression tries to **estimate these** using data.
- Size:  $d \times 1$ .

Example:

- If the true model is  $y = 2x_1 + 3x_2 + 5$ , then

$$\theta^* = [5, 2, 3]^T.$$

---

## 4. $\epsilon$ (Noise):

- Random error term captures things not explained by the model.
- Represents measurement error, randomness, or hidden factors.
- Typically assumed:

$$\epsilon \sim \mathcal{N}(0, \sigma^2 I)$$

→ Mean = 0, variance =  $\sigma^2$ , independent across data points.

Example:

- If house price is affected by "luck" (e.g., buyer really likes the house), this randomness goes into  $\epsilon$ .
- 

### 5. $\hat{\theta}$ (Estimated Coefficients):

- What regression actually computes from data.
- Formula:

$$\hat{\theta} = (X^T X)^{-1} X^T y$$


---

### 6. Error in Coefficients ( $\hat{\theta} - \theta^*$ ):

- The difference between the estimated weights and the true underlying weights.
- Can be written as:

$$\hat{\theta} - \theta^* = (X^T X)^{-1} X^T \epsilon$$

Meaning:

- The **noise in data ( $\epsilon$ ) directly affects coefficient accuracy**.
  - If features in  $X$  are highly correlated, this error gets amplified (since  $X^T X$  becomes unstable).
- 

 So in summary:

- $y$ : observed outputs
- $X$ : features
- $\theta^*$ : true model weights
- $\epsilon$ : noise/randomness
- $\hat{\theta}$ : estimated weights from regression

Noise  $\rightarrow$  propagates into  $\hat{\theta}$   $\rightarrow$  affects generalization ability.

---

### c) Probability Bounds on Error

- With random noise  $\epsilon$ \epsilon, regression error can be expressed probabilistically.

- Example (variance of coefficients):

$$\text{Var}(\hat{\theta}) = \sigma^2(X^T X)^{-1}$$

- Implications:

- More data points ( $n$ )  $\rightarrow$  smaller variance  $\rightarrow$  better generalization.
- Highly correlated features  $\rightarrow$  larger variance  $\rightarrow$  poor generalization.

## d) Error Convergence

We want to understand: as we collect **more data**, how do regression errors (in-sample and out-of-sample) behave?

### 1. Model Setup

True model:

$$y = X\theta^* + \epsilon$$

where:

- $X$  = feature matrix ( $n \times d$ )
- $\theta^*$  = true underlying coefficients
- $\epsilon \sim \mathcal{N}(0, \sigma^2 I)$  = random noise

Regression estimate:

$$\hat{\theta} = (X^T X)^{-1} X^T y$$

Prediction:

$$\hat{y} = X\hat{\theta}$$

### 2. Expected Error

We're interested in **mean squared error (MSE)**:

$$E[(y - \hat{y})^2]$$

Break it down into:

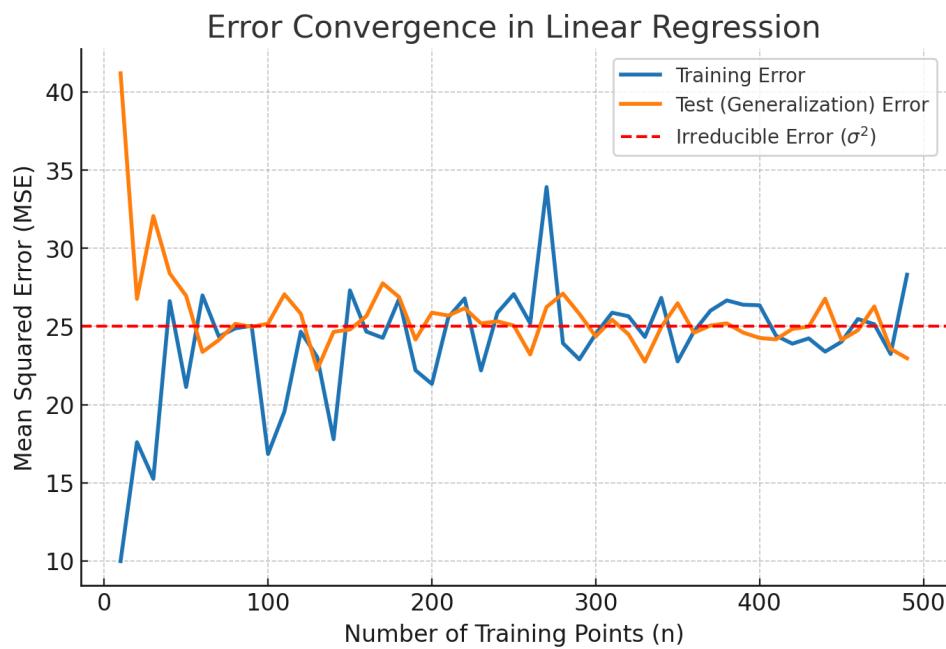
- **Bias** (systematic error from wrong model assumptions)
- **Variance** (error due to randomness in training data/noise).

In linear regression with the correct model, **bias = 0**, so error comes mainly from **variance due to noise**.

### 3. Error Convergence Intuition

- With **few data points (small n)**:
  - The estimated coefficients  $\hat{\theta}$  fluctuate a lot due to noise.
  - High **variance** in predictions  $\rightarrow$  generalization error is large.
- With **many data points (large n)**:
  - The term  $(X^T X)^{-1}$  becomes more stable.
  - The effect of random noise averages out.
  - Coefficient estimates  $\hat{\theta}$  converge to the true  $\theta^*$ .
  - Out-of-sample error approaches the irreducible noise variance  $\sigma^2$ .

Here's the **graphical example of Error Convergence**



- Training Error (blue):** starts low (since the model can "memorize" small datasets), then gradually stabilizes.

- **Test Error (orange):** initially high with few points (unstable coefficients), but decreases as  $n$  grows.
- **Irreducible Error (red dashed line):** the floor at  $\sigma^2$  — even with infinite data, test error can't go below this, because noise is unavoidable.

This illustrates that **as the sample size increases, out-of-sample error converges to the noise variance.**

---

#### 4. Mathematical Expression

The variance of estimated coefficients:

$$\text{Var}(\hat{\theta}) = \sigma^2(X^T X)^{-1}$$

- As  $n \rightarrow \infty$ :
  - $X^T X$  grows larger  $\rightarrow (X^T X)^{-1} \rightarrow 0$ .
  - $\text{Var}(\hat{\theta}) \rightarrow 0$ .
  - So,  $\hat{\theta} \rightarrow \theta^*$ .
- Out-of-sample error converges to:

$$E_{out} \rightarrow \sigma^2$$

This  $\sigma^2$  is called the **irreducible error** — the noise inherent in the system that no model can remove.

---

#### 5. Analogy

Think of regression as trying to **guess the average height of people in a city**:

- If you sample **5 people**, your estimate fluctuates wildly.
- If you sample **5000 people**, the law of large numbers ensures your estimate converges to the true population mean.
- Similarly, regression coefficients stabilize as data grows.

---

#### 6. Key Takeaways

- With small data  $\rightarrow$  error is high due to noise + unstable coefficients.

- As data increases, coefficients stabilize and error decreases.
  - Eventually, error converges to the **irreducible noise level ( $\sigma^2$ )**.
  - More data **cannot reduce error below  $\sigma^2$** , since noise is unavoidable.
  - This explains why **adding data improves generalization** (up to the noise limit).
- 

## 2. Key Takeaways

- Regression solution =  $\hat{\theta} = (X^T X)^{-1} X^T y$ .
  - Noise in data propagates into coefficient estimates.
  - Variance of coefficients depends on:
    - Noise level ( $\sigma^2$ )
    - Conditioning of feature matrix ( $X^T X$ )
  - With enough data, out-of-sample error converges to noise variance.
  - Generalization improves as:
    - $n \rightarrow \infty$  (more data)
    - Features are less correlated (well-conditioned design matrix).
- 

# Data Complexity vs Model Complexity (Guidelines for Number of Data Points vs Features)

---

## 1. Detailed Conceptual Overview

In regression (and ML broadly), the balance between the **amount of data available (data complexity)** and the **number of features used (model complexity)** is crucial.

If the model is too complex for the amount of data, it **overfits**.

If the model is too simple for the data, it **underfits**.

This topic lays out **practical guidelines** for how many data points are needed relative to the number of features.

---

### a) Dimensionality and Sample Size

- Suppose we have  $d$  features (input variables).
  - To estimate regression coefficients  $\theta$ :
    - We need at least  $d$  data points to find a unique solution.
    - In practice, we need **much more than  $d$**  for stable generalization.
- 

### b) Rules of Thumb

#### 1. Underdetermined Case ( $n < d$ ):

- More features than data points.
- Infinite possible regression solutions  $\rightarrow$  unstable coefficients.
- Overfitting is guaranteed.

#### 2. Barely Determined ( $n \approx d$ ):

- Unique solution exists (if  $X$  is full rank).
- But extremely sensitive to noise.
- Poor generalization.

#### 3. Overdetermined Case ( $n > d$ ):

- Many more data points than features.
  - Regression finds stable coefficients.
  - Best generalization.
- 

### c) Practical Guidelines in ML

- **At least 10× more data points than features** is a common heuristic for linear regression.
  - Example: If  $d = 20$  features, aim for  $n \geq 200$  samples.
- If dataset is smaller, consider:

- **Regularization** (Ridge, Lasso).
  - **Feature selection** (removing irrelevant predictors).
  - **Dimensionality reduction** (PCA, etc.).
- 

#### d) Complexity Tradeoff

- **More features (higher d):**
    - Model becomes more flexible → fits data better.
    - But risks overfitting → higher variance.
  - **More data points (higher n):**
    - Reduces variance of coefficients.
    - Improves generalization.
  - Goal = Find the right balance between data and features.
- 

## 2. Key Takeaways

- Regression stability depends on **ratio of data points (n) to features (d)**.
  - **n < d**: Underdetermined, infinite solutions, overfitting.
  - **n ≈ d**: Unique solution, but unstable due to noise.
  - **n >> d**: Stable, generalizable regression.
  - **Rule of thumb**:  $n \geq 10d$  for reliable regression.
  - If data is limited → use **regularization, feature selection, or dimensionality reduction** to control complexity.
- 

# Statistical Inference in Regression (t-statistic and p-value)

---

## 1. Detailed Conceptual Overview

So far, regression has given us **coefficients** ( $\hat{\theta}$ ) that describe the relationship between features and the target. But how do we know **which coefficients are truly meaningful?**

This is where **statistical inference** comes in — it helps us move from *just fitting a model* to *evaluating whether the model's parameters are significant or just noise*.

The two key tools here are:

- **t-statistic:** Measures how strongly a coefficient differs from zero (relative to its variability).
- **p-value:** Quantifies the probability that the observed effect could have occurred by random chance.

Together, these allow us to test hypotheses about the importance of features.

---

### a) Standard Error of Coefficients

- When we estimate a regression coefficient  $\hat{\theta}_j$ , it is not exact.
- Different samples (drawn from the same population) would give slightly different  $\hat{\theta}_j$ .
- The **Standard Error (SE)** tells us how much  $\hat{\theta}_j$  tends to vary across repeated samples.
- Think of it as the “uncertainty” or “spread” of our estimated coefficient.

So:

- **Small SE** → coefficient estimate is stable, reliable.
- **Large SE** → coefficient estimate is noisy, less trustworthy.

Mathematically:

$$SE(\hat{\theta}_j) = \sqrt{\sigma^2 \cdot [(X^T X)^{-1}]_{jj}}$$

where:

- $\sigma^2$  = estimated variance of noise,
- $(X^T X)^{-1}_{jj}$  = diagonal element corresponding to feature j.

## Where Does SE Come From?

In linear regression:

$$\hat{\theta} = (X^T X)^{-1} X^T y$$

If the true data-generating process is

$$y = X\theta^* + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2 I),$$

then the variance of the estimated coefficients is:

$$\text{Var}(\hat{\theta}) = \sigma^2 (X^T X)^{-1}$$

---

### b) Hypothesis Testing for Coefficients

For each coefficient  $\theta_j$ :

- **Null hypothesis**  $H_0: \theta_j = 0$  (feature has no effect).
- **Alternative hypothesis**  $H_1: \theta_j \neq 0$  (feature is important).

We want to know if the data provides enough evidence to reject  $H_0$ .

---

### c) t-Statistic

The **t-statistic** is defined as:

$$t_j = \frac{\hat{\theta}_j}{SE(\hat{\theta}_j)}$$

Interpretation:

- A large  $t_j$  (positive or negative) means the coefficient is many standard errors away from zero  $\rightarrow$  strong evidence against  $H_0$ .
  - A small  $t_j$  means the coefficient could plausibly be zero.
- 

### d) p-Value

The **p-value** is a probability. Specifically, it answers this question:

If the null hypothesis is true (coefficient = 0, no effect), how likely is it that we would observe a coefficient estimate as extreme (or more extreme) than the one we got?

So:

- **Small p-value** → It would be very unlikely to see such a strong coefficient if the feature truly had no effect → Evidence *against* the null → Feature is significant.
- **Large p-value** → The observed effect could easily happen by chance noise → No evidence that the feature matters.
- If  $p < \alpha$  (commonly  $\alpha = 0.05$ ): Reject  $H_0$ , feature is statistically significant.
- If  $p \geq \alpha$ : Fail to reject  $H_0$ , the feature is not significant.

Example interpretations:

- $p = 0.001$ : Strong evidence that the feature matters.
- $p = 0.3$ : Weak evidence, feature likely unimportant.

## Common Misconceptions

### X “A small p-value means the effect is large.”

- Wrong. It means the effect is unlikely to be zero, but it doesn't measure *how big* it is.

### X “A p-value tells us the probability that the null hypothesis is true.”

- Wrong. The p-value is computed **assuming the null is true**. It tells us how surprising the data would be under that assumption.

### ✓ Correct interpretation:

- The p-value measures the compatibility of the data with the null hypothesis.

## e) Why We Need This

- **Not all coefficients matter:** Some may appear large due to noise, but not truly significant.
- **Helps with feature selection:** By checking significance, we can identify which predictors add real value.
- **Avoids over-interpretation:** Prevents us from drawing conclusions based on random fluctuations in data.

### f) Example Scenario

Imagine a regression predicting **Savings** from features: Salary, Age, Education, Inherited Wealth.

- **Salary:** Coefficient = 0.8, SE = 0.1  $\rightarrow t = 8$ ,  $p < 0.001 \rightarrow$  Highly significant.
- **Education:** Coefficient = 0.05, SE = 0.08  $\rightarrow t = 0.62$ ,  $p = 0.54 \rightarrow$  Not significant.
- **Inherited Wealth:** Coefficient = 0.4, SE = 0.15  $\rightarrow t = 2.67$ ,  $p = 0.008 \rightarrow$  Significant.

From this, we conclude: Salary and Inherited Wealth are strong predictors, while Education has little effect after accounting for others.

---

## 2. Key Takeaways

- Regression coefficients come with **uncertainty**, captured by **standard errors**.
  - **t-statistic** compares the coefficient size relative to its variability.
  - **p-value** quantifies the probability of seeing such a coefficient if the true effect were zero.
  - **Small p-values ( $< 0.05$ )**  $\rightarrow$  feature is significant.
  - **Large p-values**  $\rightarrow$  feature likely not important.
  - Statistical inference helps us distinguish **real predictors from noise**, guiding better model interpretation and feature selection.
-