

# Naïve Bayes

---



## Intuition: What is Naïve Bayes?

Naïve Bayes is a **probabilistic classifier** based on **Bayes' Theorem** — but with a “naïve” assumption of independence between features.

### Bayes' Theorem (the foundation)

$$P(y|x) = \frac{P(x|y), P(y)}{P(x)}$$

where:

- ( $y$ ) = class label
- ( $x = (x_1, x_2, \dots, x_n)$ ) = feature vector
- ( $P(y|x)$ ) = probability of class given features (**posterior**)
- ( $P(x|y)$ ) = likelihood of features given class
- ( $P(y)$ ) = prior probability of class
- ( $P(x)$ ) = probability of features (normalization constant)



### The “Naïve” Assumption

The **naïve** part assumes that *all features are conditionally independent* given the class label.

Mathematically:

$$P(x_1, x_2, \dots, x_n|y) = \prod_{i=1}^n P(x_i|y)$$

This drastically simplifies the computation of ( $P(x|y)$ ).

So the model becomes:

$$P(y|x_1, \dots, x_n) \propto P(y) \prod_{i=1}^n P(x_i|y)$$



In other words:

The “naïve” assumption says:

“Let’s pretend that all features are independent of each other — once we know the class.”

In math:

$$P(x_1, x_2, \dots, x_n | y) = P(x_1 | y) \times P(x_2 | y) \times \dots \times P(x_n | y)$$

That’s it — this is the **Naïve Bayes assumption**.

---

## In Plain English

Let’s say you’re predicting whether an email is **spam**.

You have features like:

- $x_1$  = contains “free”
- $x_2$  = contains “win”
- $x_3$  = contains “meeting”

The **naïve assumption** says:

Once we know the email is spam, whether it contains “free” has nothing to do with whether it contains “win” or “meeting.”

So instead of learning a big complex joint probability like  $P(\text{free}, \text{win}, \text{meeting} | \text{spam})P(\text{free}, \text{win}, \text{meeting} | \text{spam})P(\text{free}, \text{win}, \text{meeting} | \text{spam})$ ,

We just multiply the separate probabilities:

$$P(\text{free} | \text{spam}) \times P(\text{win} | \text{spam}) \times P(\text{meeting} | \text{spam})$$

---

## Why Make This Assumption?

Because it makes everything **computationally simple**.

- Without the assumption: we’d need exponential data to estimate all combinations of features.
- With the assumption: we only need a few probabilities per feature and class.

Even though the assumption is *not really true* (features do correlate),

Naïve Bayes still performs **surprisingly well** — especially for text data, where independence is *approximately true*.

---



## Prediction Rule

We predict the class ( $y$ ) that maximizes the posterior probability:

$$\hat{y} = \arg \max_y, P(y) \prod_{i=1}^n P(x_i|y)$$

So it's just about computing **simple probabilities**, multiplying them, and choosing the class with the highest result.

---



## Types of Naïve Bayes

Type	Description	Example Use
<b>Gaussian Naïve Bayes</b>	Assumes continuous features are distributed normally within each class	Continuous numeric data (e.g., sensor readings)
<b>Multinomial Naïve Bayes</b>	Used for <b>count features</b> like word frequencies	Text classification, spam filtering
<b>Bernoulli Naïve Bayes</b>	Used for <b>binary features</b> (0/1, yes/no)	Sentiment analysis, document classification with presence/absence features
<b>Categorical Naïve Bayes</b>	Works with discrete categorical variables	Customer segmentation, survey data

---



## Use Cases

Naïve Bayes shines in:

- **Spam detection** (words like "free", "winner" → spam class)
- **Text classification** (sentiment analysis, topic tagging)
- **Medical diagnosis** (symptom presence → disease class)
- **Document categorization** (news, tweets, reviews)
- **Recommender systems** (predicting user preferences)

Essentially, it's best when:

- You have **lots of features** (like words in a vocabulary)
  - Each feature gives **small independent evidence** about the class
- 

## ✓ Benefits

Benefit	Explanation
<b>Simple and fast</b>	Only needs frequency counts or Gaussian parameters — no iterative training
<b>Works well with small data</b>	Even limited training data can give good estimates
<b>Scalable</b>	Linear in number of features and samples
<b>Performs surprisingly well in practice</b>	Despite the "naïve" assumption, often competitive with complex models (especially in text tasks)
<b>Requires little storage</b>	Just needs to store probabilities $P(x_i y)$
<b>Handles high-dimensional data well</b>	Works great for text where features (words) are numerous and sparse
<b>Probabilistic output</b>	Gives interpretable probabilities for each class

---

## ⚠ Limitations

Limitation	Why
<b>Strong independence assumption</b>	Real-world features often correlate (e.g., "expensive" and "luxury" words)
<b>Zero-frequency problem</b>	If a feature never occurs in training for a class, its probability becomes zero → fixed by <b>Laplace smoothing</b>
<b>Not ideal for correlated or continuous data</b>	Works best when features truly are conditionally independent

---

## 🔍 Example: Spam Detection (Simplified)

Let's say you have two classes:

## **Spam ( $y=1$ ) and Not Spam ( $y=0$ )**

And features:

$x_1$  = "contains free",  $x_2$  = "contains win",  $x_3$  = "contains meeting"

Naïve Bayes would compute:

$$P(\text{spam}|x) \propto P(\text{spam}) \times P(\text{free}|\text{spam}) \times P(\text{win}|\text{spam}) \times P(\text{meeting}|\text{spam})$$

and similarly for "Not Spam," then choose whichever is larger.

---

## **Summary**

Concept	Naïve Bayes Summary
<b>Type</b>	Generative probabilistic classifier
<b>Core idea</b>	Use Bayes' theorem + assume feature independence
<b>Equation</b>	$P(y) \prod_{i=1}^n P(x_i y)$
<b>Output</b>	Class with highest posterior probability
<b>Strengths</b>	Fast, simple, robust, effective for text
<b>Weaknesses</b>	Assumes independence, zero-frequency issue
<b>Common Variants</b>	Gaussian, Multinomial, Bernoulli

---