

Nearest Neighbor methods

KNN Regression

Definition

K-Nearest Neighbors (KNN) Regression is a **non-parametric supervised learning algorithm** that predicts the value of a target variable for a new data point by looking at the values of its **K nearest neighbors** in the training data. Instead of predicting a class (as in KNN classification), it predicts a **numerical value**.

How It Works

1. Choose K

- Decide the number of neighbors K.
- Example: K = 3.

2. Distance Metric

- Compute the distance (commonly **Euclidean distance**) between the new data point and all training points.

$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

3. Find Nearest Neighbors

- Select the K training points that are closest to the new data point.

4. Aggregate Neighbor Values

- For regression, take the **average (mean)** of the target values of the K neighbors:

$$\hat{y} = \frac{1}{K} \sum_{i=1}^K y_i$$

- Sometimes, a **weighted average** is used, giving closer neighbors more influence:

$$\hat{y} = \frac{\sum_{i=1}^K \frac{1}{d_i} y_i}{\sum_{i=1}^K \frac{1}{d_i}}$$

Example

Suppose you want to predict a house price based on **square footage**.

- Training data:
 - 1000 sqft → \$200k
 - 1500 sqft → \$250k
 - 1700 sqft → \$270k
 - 3000 sqft → \$500k
- Query: 1600 sqft, K=2.
- Nearest neighbors: 1500 sqft (\$250k), 1700 sqft (\$270k).
- Predicted price:

$$\hat{y} = \frac{250k + 270k}{2} = 260k$$

Advantages

- **Simple & intuitive** – easy to implement.
- **No training step** – just store data.
- **Flexible** – can model non-linear relationships.

Disadvantages

- **Computationally expensive** – needs distance calculations for all points at prediction time.
 - **Sensitive to choice of K.**
 - **Sensitive to irrelevant features and scale of data** → usually requires normalization/standardization.
-

Choosing K in KNN Regression

1. Small K (e.g., 1, 2, 3)

- **Pros:** Model is very flexible, captures local patterns well.
 - **Cons:** Predictions become noisy (high variance), sensitive to outliers.
 - Example: If K=1, you just copy the nearest neighbor's value (can be unstable).
-

2. Large K (e.g., 20, 50, 100)

- **Pros:** Predictions are smoother, less affected by noise.
 - **Cons:** Model becomes too rigid (high bias), may miss important local variations.
 - Example: If K is too large, you're basically averaging over the whole dataset → prediction tends toward the global mean.
-

3. Rule of Thumb

- Start with:

$$K \approx \sqrt{N}$$

where N = number of training samples.

- Example: If you have 100 training points, start testing with $K \approx 10$.
-

4. Cross-Validation (Best Method)

- Use **k-fold cross-validation**:
 1. Try different K values (e.g., 1 → 30).

2. Measure error (e.g., Mean Squared Error, RMSE).
 3. Pick the K that minimizes error on validation sets.
-

5. Weighted KNN

- Instead of worrying too much about the exact K, you can assign **weights based on distance**:
 - Closer neighbors get a higher weight.
 - This reduces the problem of choosing a "perfect" K.
-

Visual Intuition

- **Small K** → "wiggly" prediction curve, overfitting.
 - **Large K** → very smooth, underfitting.
 - **Optimal K** → balance between bias and variance.
-

In practice:

- Try several K values using **cross-validation**.
 - Prefer **odd K** in classification (to avoid ties), but in regression, odd/even doesn't matter.
 - Normalize features first (since distance is sensitive to scale).
-

KNN Regression VS Linear Regression

◆ 1. Nature of the Model

- **KNN Regression:**
 - **Non-parametric** (doesn't assume any specific function form).
 - Prediction is based on **local neighbors**.
- **Linear Regression:**
 - **Parametric** (assumes a linear relationship between features and target).

- Fits a global line/plane through all the data.
-

◆ 2. How Prediction Works

- **KNN Regression:**

- For a new point, look at the **K nearest neighbors** and average their target values.
- Local approach → depends on nearby data.

- **Linear Regression:**

- Learns coefficients β for each feature using training data.
- Prediction uses:

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_n x_n$$

◆ 3. Training vs Testing

- **KNN Regression:**

- No real *training phase*. Just store data.
- **Slow at prediction** (needs distance calculation each time).

- **Linear Regression:**

- Needs a **training phase** (fit coefficients).
 - **Fast at prediction** (just plug into equation).
-

◆ 4. Interpretability

- **KNN Regression:** Harder to interpret, just averages neighbors.
 - **Linear Regression:** Very interpretable (coefficients tell how much each feature affects target).
-

◆ 5. Handling Relationships

- **KNN Regression:**

- Can capture **non-linear relationships**.
 - Works well if patterns are local.
 - **Linear Regression:**
 - Only models **linear relationships** (unless you add polynomial features or interactions).
-

◆ 6. Sensitivity

- **KNN Regression:**
 - Sensitive to irrelevant features and scaling (always normalize data!).
 - Sensitive to choice of K.
 - **Linear Regression:**
 - Sensitive to outliers (one extreme point can distort the line).
 - Assumes independence, homoscedasticity, and no multicollinearity.
-

◆ 7. Example: Predicting House Price by Square Footage

- Suppose houses are around a curve-shaped trend (non-linear).
- 👉 **KNN Regression**
- Looks at nearby houses, averages their prices → adapts to local curve.
- 👉 **Linear Regression**
- Fits a straight line. If the relationship is not linear, predictions will be biased.
-

◆ Summary Table

Aspect	KNN Regression 	Linear Regression 
Model type	Non-parametric	Parametric
Training phase	None	Required (fit line)
Prediction speed	Slow	Fast
Relationship handled	Non-linear ok	Only linear

Aspect	KNN Regression 	Linear Regression 
Interpretability	Low	High
Sensitive to	Scale, K, neighbors	Outliers, assumptions

 In short:

- Use **Linear Regression** if the relationship is approximately linear and interpretability matters.
- Use **KNN Regression** if the relationship is unknown/non-linear and you have enough data.

◆ Why KNN Works Well for Non-Linear Data

- **Linear Regression** assumes:

$$y = \beta_0 + \beta_1 x_1 + \cdots + \beta_n x_n$$

→ a straight line (or hyperplane). If the data has a **curved relationship**, it struggles unless we add polynomial terms or transformations.

- **KNN Regression** makes **no assumption** about the data's shape.
 - It predicts by looking at nearby points and averaging their values.
 - This means it naturally adapts to **curves, clusters, or any non-linear patterns**.

◆ Example: Non-Linear Relationship

Suppose the true relationship is:

$$y = \sin(x) + \text{noise}$$

- **Linear Regression**: Will try to fit a straight line → poor fit.
- **KNN Regression**:
 - At each x, it looks around locally and averages nearby y values.
 - The prediction curve bends and follows the sine wave shape.

◆ Visual Intuition

Imagine plotting the data points shaped like a **U-curve (quadratic)**:

- **Linear Regression:**

- Fits a straight diagonal line through the points.
- High bias, misses the "U" shape.

- **KNN Regression:**

- Each prediction is based on local neighbors.
- Predictions curve along the data naturally, following the "U".

◆ Trade-off (Bias-Variance)

- **Small K** → very wiggly curve, may overfit noise.
- **Large K** → very smooth, may underfit (looks too linear).
- Cross-validation helps find the sweet spot.

✓ Key takeaway:

KNN is especially powerful for **non-linear patterns** because it doesn't assume any global form; it adapts locally to the data.



KNN works better for lower numbers of features and more nonlinear data (for $d \leq 3$)

For linear relationships and higher-dimensional data, linear regression is generally more effective and provides a better generalized model (for $d \geq 4$)

KNN Classification:K-Nearest Neighbors (KNN)

◆ What is KNN?

K-Nearest Neighbors (KNN) is a **supervised machine learning algorithm** used for **classification** and **regression**.

- It works by **comparing new data to stored training data**.
- The prediction is made based on the **K closest (nearest) neighbors** in feature space.

It's called **lazy learning** because:

- No model is built during training.
- The algorithm just stores the dataset.
- Computation happens only when making a prediction.

◆ How KNN Works (Step by Step)

1. **Choose K** → the number of neighbors to consider.
 2. **Compute Distance** → between the query point and all training points.
 - Common metric: **Euclidean distance**
$$d(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$
 3. **Select Nearest Neighbors** → pick the K closest points.
 4. **Make Prediction:**
 - **Classification** → majority vote of neighbor labels.
 - **Regression** → average (or weighted average) of neighbor target values.
-

◆ KNN for Classification

- Output = **class label** (e.g., "cat" vs "dog").
- Prediction based on **majority voting** among neighbors.

Example: K=5, neighbors = {Dog, Dog, Cat, Dog, Cat} → Prediction = **Dog**.

◆ KNN for Regression

- Output = **numerical value**.
- Prediction = **mean (or weighted mean)** of neighbor values.

Example: Predict house price → average prices of 5 nearest houses.

◆ Choosing K

- **Small K (e.g., 1, 3)** → sensitive to noise (high variance).
 - **Large K (e.g., 20, 50)** → smoother predictions but may underfit (high bias).
 - Best K is chosen using **cross-validation**.
 - Rule of thumb: $K \approx \sqrt{N}$, where N = number of samples.
-

◆ Pros and Cons

✓ Advantages

- Simple and intuitive.
- Works for both classification and regression.
- Non-parametric → no assumption about data distribution.
- Can model **non-linear** decision boundaries.

✗ Disadvantages

- Computationally expensive at prediction time.
 - Sensitive to feature scaling (need normalization).
 - Struggles in high dimensions (curse of dimensionality).
 - Sensitive to noisy/irrelevant features.
-

◆ Example (Classification)

Dataset: Predict whether fruit is Apple  or Orange  based on **weight** and **color**.

- Training:
 - (150g, Red) → Apple

- (180g, Red) → Apple
 - (120g, Orange) → Orange
 - (130g, Orange) → Orange
 - Query: (160g, reddish).
 - Distances → nearest neighbors = 2 Apples, 1 Orange → Prediction = **Apple**.
-

 **In short:**

KNN is a **memory-based algorithm** that predicts using **neighbors**:

- **Classification** = majority vote
 - **Regression** = average value
-

Decision Boundary

Definition

A **decision boundary** is the line (in 2D), curve, or surface (in higher dimensions) that separates different **class regions** in the feature space, based on the predictions of a model.

- On one side of the boundary, → model predicts **Class A**
- On the other side → model predicts **Class B**

It's essentially the "border" where the model is undecided.

◆ Example in 2D

Imagine a dataset with two features:

- x_1 = weight of fruit
- x_2 = color intensity

And two classes: **Apple (red)**, **Orange (orange)**.

- The **decision boundary** is the curve/line where the model is equally likely to predict Apple or Orange.

- If you plot the data points in 2D space, the boundary "cuts" the plane into regions → each region belongs to one class.
-

◆ Decision Boundaries by Algorithm

1. Linear Models (Logistic Regression, Linear SVM)

- Decision boundary is a **straight line (2D)** or **hyperplane (nD)**.

- Equation looks like:

$$w_1x_1 + w_2x_2 + b = 0$$

- Example: Separate apples and oranges with a straight line.
-

1. KNN (K-Nearest Neighbors)

- The decision boundary is **non-linear and jagged**.
 - It depends on the local arrangement of points.
 - For K=1, every training point creates its own "region" (Voronoi diagram).
 - For larger K, the boundaries are smoother.
-

1. Decision Trees

- Boundaries are **axis-aligned** (vertical or horizontal splits in 2D).
 - Looks like a series of rectangles partitioning the feature space.
-

1. Non-linear Models (Kernel SVM, Neural Nets, KNN)

- Decision boundaries can be **curved and complex**.
 - Adapt better to non-linear patterns in the data.
-

◆ Why It Matters

- The decision boundary shows **how the model generalizes** to unseen data.
 - **Too complex a boundary** = overfitting (follows noise).
 - A **too simple boundary** = underfitting (misses real patterns).
-

◆ Visualization Intuition

- Imagine a map where apples  and oranges  are scattered.
 - The **decision boundary** is the “fence” the algorithm builds between the apple-region and the orange-region.
 - If a new fruit falls inside the apple region, the model says **Apple**.
-

✓ In short:

The **decision boundary** is the dividing line/surface in feature space that tells us which class a new point belongs to. Its **shape** depends on the algorithm:

- Straight for linear models
 - Jagged for KNN
 - Step-like for trees
 - Curvy for non-linear methods
-

Error Rate in KNN

Definition

The **error rate** in KNN is the fraction of predictions the algorithm gets **wrong** compared to the true labels.

Mathematically:

$$\text{Error Rate} = \frac{\text{Number of Incorrect Predictions}}{\text{Total Number of Predictions}}$$

Equivalently:

Accuracy=1–Error Rate

◆ How Error Rate is Measured

1. Training Error (Resubstitution Error)

- Run KNN on the training dataset itself.

- With $K=1$, training error is **0** (every point is its own nearest neighbor).
- But this doesn't mean good generalization → likely overfitting.

2. Test Error

- Run KNN on a separate test set.
- Shows how well the model generalizes to unseen data.

3. Cross-Validation Error

- Split data into k folds.
 - Train on $k-1$ folds, test on the remaining fold, repeat.
 - Average error across folds → more reliable estimate.
-

◆ Error Rate vs K (Bias-Variance Tradeoff)

- **Small K (e.g., 1)**
 - Training error ≈ 0 (perfect memorization).
 - Test error is high → overfitting, high variance.
- **Large K (e.g., 50)**
 - Model becomes too smooth.
 - Training and test error both increase → underfitting, high bias.
- **Optimal K**
 - Somewhere in between → lowest test error.
 - Found by cross-validation.

 If you plot **Error Rate vs K**:

- Training error starts very low, increases with K .
 - Test error decreases at first, then increases again (U-shape curve).
-

◆ Example

Suppose we have 100 test samples, and KNN predicts 85 correctly.

- Correct predictions = 85
- Incorrect predictions = 15

$$\text{Error Rate} = \frac{15}{100} = 0.15 = 15\%$$

$$Accuracy = 85$$

 **In short:**

- **Error rate in KNN** = fraction of wrong predictions.
 - Depends heavily on **choice of K** and **data distribution**.
 - Best K is chosen by minimizing **cross-validation error**.
-

◆ Variants of KNN

1. Weighted KNN

- Standard KNN: all neighbors contribute equally.
- Weighted KNN: closer neighbors have **higher influence**.
- Prediction:

$$\hat{y} = \frac{\sum_{i=1}^K w_i y_i}{\sum_{i=1}^K w_i}, \quad w_i = \frac{1}{d_i}$$

- Useful when local points are much more relevant than distant ones.
-

2. Distance-Weighted KNN with Kernel Functions

- Instead of just $1/d$, use kernel functions to weight neighbors.
 - Examples:
 - Gaussian kernel
 - Epanechnikov kernel
 - Smooths predictions better.
-

3. Condensed KNN (cKNN)

- Stores only a **subset of representative points** (prototypes).
 - Reduces memory usage and speeds up prediction.
-

4. Reduced KNN (rKNN)

- Removes training samples that don't affect decision boundaries.
 - Smaller, faster model without losing accuracy.
-

5. Edited KNN (eKNN)

- Cleans noisy data points by removing mislabeled/irregular samples.
 - Improves robustness against outliers.
-

6. Fuzzy KNN

- Instead of a hard decision (class A or B), assigns **membership probabilities**.
 - Example: "70% Cat, 30% Dog" instead of just "Cat".
-

7. KNN with Dimension Reduction

- Apply **PCA, LDA, t-SNE, UMAP** before KNN to fight the **curse of dimensionality**.
 - Improves performance in high-dimensional datasets (e.g., text, images).
-

8. Approximate Nearest Neighbors (ANN)

- Exact KNN is slow for large datasets.
 - ANN methods (KD-Trees, Ball Trees, Locality Sensitive Hashing) speed up neighbor search.
 - Popular in recommender systems and similarity search.
-

9. Adaptive KNN

- Instead of a fixed K, adaptively choose K based on local density.
 - In dense areas → small K.
 - In sparse areas → larger K.
-

10. Parallel / Scalable KNN

- Distributed KNN (e.g., using Hadoop/Spark).
 - Designed for big data where classic KNN is too slow.
-

◆ Summary

Variant	Purpose
Weighted KNN	Give closer neighbors more influence
Kernel KNN	Smooth weighting with kernels
Condensed/Reduced/Edited KNN	Reduce dataset size & noise
Fuzzy KNN	Soft probability-based classification
Dimensionality reduction + KNN	Handle high-dimensional data
Approximate KNN	Fast lookup in large datasets
Adaptive KNN	Flexible K depending on density
Parallel/Scalable KNN	Works with very large datasets

✓ **In short:** The basic KNN idea is simple, but variants exist to tackle **efficiency, noise, interpretability, high-dimensional data, and scalability**.
