

# Image processing

---

## 1. Feature Engineering for Image Data

### ◆ What It Means

Feature engineering in the context of image data involves converting raw pixel values into a form that machine learning algorithms can use effectively — typically vectors or tensors.

### ◆ Why It's Needed

Machine learning models cannot process raw image files (like `.jpg`, `.png`) directly. Instead, these images must be **numerically encoded** to represent spatial and visual information (colors, shapes, textures, etc.).

### ◆ Key Concepts

- **Image → Tensor:** An image is transformed into a multidimensional array (tensor) — e.g., a color image becomes a 3D array (height × width × channels).
- **Feature Vector:** You can flatten or reduce these tensors into 1D arrays for algorithms like k-means or SVM, though CNNs work directly with tensors.

---

## 2. Handling Colour Images

### ◆ RGB Color Channels

Color images use the **RGB color model**, which includes:

- **Red Channel**
- **Green Channel**
- **Blue Channel**

Each channel is a **2D matrix** of pixel intensity values between 0 and 255 (for 8-bit images).

### ◆ Image Shape

For example, an image of size 32×32 will have:

Shape: (32, 32, 3)

Where:

- 32 is the height
- 32 is the width
- 3 is the number of color channels (R, G, B)

## ◆ Grayscale Conversion

Sometimes color images are converted to **grayscale** by averaging or using weighted sums of the RGB channels to simplify the feature extraction process.

## 🔍 3. Incorporating Local Context Using Convolutions

### ◆ What is a Convolution?

A **convolution** is a mathematical operation where a small matrix called a **filter** or **kernel** slides across the image to compute **dot products** with local patches.

For a 3×3 filter applied to a 5×5 image patch:

```
[1 0 -1]   [img11 img12 img13]
[1 0 -1] * [img21 img22 img23]
[1 0 -1]   [img31 img32 img33]
```

### ◆ Why Use Convolutions?

- Preserve **spatial locality** — neighboring pixels often share semantic meaning.
- Extract **meaningful patterns** — edges, corners, textures.
- Mimic **human visual perception** where localized features are crucial.

### ◆ Filter Types

Filters (kernels) can be:

- **Smoothing** (e.g., mean filter: all ones divided by size)
- **Edge detection** (e.g., Sobel, Prewitt)

- **Sharpening** (emphasize borders)

## ◆ Output

After applying a filter:

- You get a **new image** (convolved feature map).
- Each value represents a **local pattern match** between the filter and the input image patch.

## ◆ Stacking Feature Maps



You can apply multiple filters to get a **stack of feature maps** — each highlighting a different type of pattern. These form the input to deeper layers in CNNs or can be flattened into feature vectors for classical ML.

# 1. Basic Methods of Image Feature Extraction

## A. Raw Pixel Intensities

- **Approach:** Flatten the 2D (grayscale) or 3D (RGB) image into a 1D vector.
- **Example:**
  - Grayscale image  $28 \times 28 \rightarrow$  vector of size 784
  - RGB image  $32 \times 32 \times 3 \rightarrow$  vector of size 3072

```
import numpy as np
image = ... # a 32×32×3 image
features = image.flatten() # shape becomes (3072,)
```

-  **Pros:** Simple and quick
-  **Cons:** Ignores spatial relationships and patterns

## B. Statistical Features

- **Histogram of pixel values:** Captures intensity distribution.
- **Mean, Variance, Skewness** of each channel.
- Often used in medical or texture-based image analysis.

```
mean = np.mean(image, axis=(0, 1))
std = np.std(image, axis=(0, 1))
```

## C. Convolution-based Features (Manual Filters)

- Use 3×3, 5×5, etc., **kernels** to extract local features like edges, corners, textures.
- Apply convolution → flatten the result → use as features.

```
from scipy.signal import convolve2d
```

```
kernel = np.ones((3, 3)) / 9 # simple blur filter
convolved = convolve2d(gray_image, kernel, mode='valid')
features = convolved.flatten()
```

- This simulates part of what CNNs do automatically.

## 2. Intermediate Techniques

### A. PCA (Principal Component Analysis)

- Reduces dimensionality while retaining variance.
- Converts flattened image vectors into a lower-dimensional space.

```
from sklearn.decomposition import PCA
```

```
pca = PCA(n_components=50)
features = pca.fit_transform(flat_images)
```

### B. Histogram of Oriented Gradients (HOG)

- Captures edge orientations.
- Common in object detection (e.g., pedestrian detection).

```
from skimage.feature import hog
```



```
features = hog(image, pixels_per_cell=(8, 8), cells_per_block=(2, 2), multic  
hannel=True)
```

## 3. Advanced Methods

### A. Pre-trained CNN Feature Extractors

- Use models like **ResNet**, **VGG**, or **MobileNet** pretrained on ImageNet.
- Remove the classification layer and extract from a hidden layer.

```
from tensorflow.keras.applications import VGG16  
from tensorflow.keras.applications.vgg16 import preprocess_input  
from tensorflow.keras.models import Model  
  
model = VGG16(weights='imagenet', include_top=False, input_shape=(224,  
224,3))  
feature_extractor = Model(inputs=model.input, outputs=model.output)  
  
features = feature_extractor.predict(preprocess_input(image_batch))  
flattened_features = features.reshape(features.shape[0], -1)
```

-  Captures very rich hierarchical features
-  Highly effective for transfer learning and clustering

### B. Autoencoders

- Train an encoder-decoder network to reconstruct the image.
- Use the bottleneck layer as the feature vector.

```
# encoder_output = encoder.predict(image_batch)
```

## Best Practices

- **Normalize pixel values:** Scale between 0 and 1 or use standardization (mean=0, std=1)
- **Resize images:** Ensure consistent shape across dataset

- **Augmentation:** For better generalization, augment with flips, rotations, etc.

## Summary Table

Method	Feature Type	Complexity	Captures Spatial Info
Raw pixels	Flat vector	Low	✗ No
Manual convolutions	Local patterns	Low-Medium	✓ Some
PCA	Reduced linear components	Medium	✗ Mostly global
HOG	Edge orientations	Medium	✓ Yes
Pre-trained CNNs	Deep hierarchical features	High	✓ Strong
Autoencoders	Latent representations	High	✓ Strong

## 1. Downsizing Images

### What It Means

**Downsizing an image** refers to **reducing its resolution** — i.e., making it smaller in terms of width and height (fewer pixels).

### Why It's Done

- **Reduce computational cost:** Smaller images require less memory and processing.
- **Speed up training:** Especially in deep learning where input size significantly affects performance.
- **Standardize input:** Most ML models and CNN architectures require fixed-size input (e.g., 224×224 for VGG).

### Example

Original size: 256×256

Downsized to: 64×64

Each image goes from 65,536 pixels → 4,096 pixels.

## ◆ How It's Done

Using `cv2` (OpenCV) or `PIL`:

```
import cv2
resized = cv2.resize(image, (64, 64))
```

## ◆ Trade-offs

- **! Information loss:** Important details like edges or fine textures may be lost.
- **✓** Best to keep enough resolution to preserve features relevant to the task.

# 2. Downsizing Color (Reducing Color Depth or Channels)

## ◆ What It Means

There are **two main interpretations**:

### A. Reducing Color Depth (Bit Depth)

- Reducing the number of bits used to represent color intensity.
- From 8 bits/channel (256 values) → 4 bits/channel (16 values), etc.

This reduces file size and can simplify features but may lead to visible color banding.

### B. Converting Color Images to Grayscale

- Reducing from **3 channels (RGB)** to **1 channel (Grayscale)**.
- Common when color information isn't essential.

```
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
```

- RGB image shape: (32, 32, 3)
- Grayscale image shape: (32, 32)

## ◆ Why It's Done

- **Simplifies models:** Less input data.
- **Speeds up processing.**

- **Avoids unnecessary complexity** if color isn't needed (e.g., digit recognition).

### ◆ Trade-offs

- **!** Color-based features are lost, which might be crucial in tasks like traffic light detection or object classification in natural images.

### ✓ Summary

Aspect	Downsizing Images	Downsizing Color
What it reduces	Spatial resolution	Color information
Input dimension	From $H \times W$ to smaller $H \times W$	From 3 channels $\rightarrow$ 1 channel
Use case	Efficiency, standardization	When color is non-essential
Trade-off	Blurry/loss of detail	Loss of color cues